



CS 564: Database Management Systems

Lecture 17: B+ Tree

Xiangyao Yu
3/1/2024

Module B2 Indexes

Hash index

B+ tree index

Advanced Indexing

External sort

Outline

B+ tree data structures

B+ tree operations

- Search
- Insertion
- Deletion

Primary vs. secondary indexes

Motivation

We have the following SQL query:

```
SELECT *  
FROM Sales  
WHERE price > 100 ;
```

Hash index accelerates **equality** search

- In expectation constant I/O cost for search and insert

Tree-based index accelerates both **equality** and **range** search

- In expectation $\text{Log}(N)$ I/O cost for search and insert

B+ Tree Index

A dynamic tree-structured index

- Adjusted to be always height-balanced
- 1 node = 1 physical page

Supports efficient **equality** and **range** search

Widely used in many DBMSs

Outline


B+ tree data structures

B+ tree operations

- Search
- Insertion
- Deletion

Primary vs. secondary indexes

B+ Tree Index – Basic Structure



Index entries:

- Exist *only* in the leaf nodes
- Are sorted according to the search key

B+ Tree Index – Node

The parameter d is the *order* of the tree


Each node contains $d \leq m \leq 2d$ entries

– Minimum 50% occupancy

With the exception of the root node, which can have $1 \leq m \leq 2d$ entries

Non-Leaf Nodes


A non-leaf (or internal) node with m entries has $m+1$ pointers to lower-level nodes



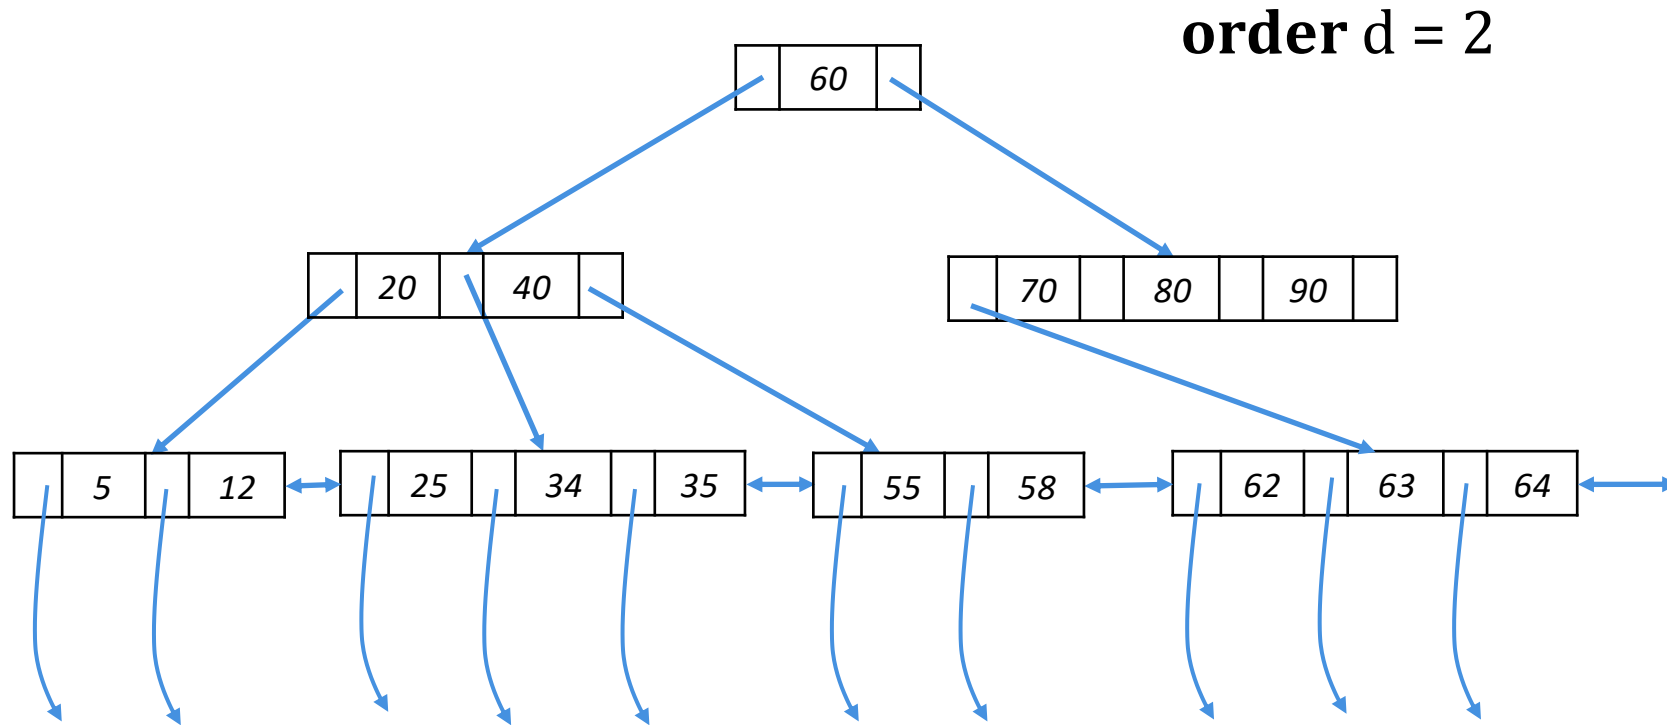
Leaf Nodes

A leaf node with m entries has



- m pointers to the data records (rids)
- pointers to the **next** and **previous** leaves



B+ Tree Example



B+ Tree vs. B-Tree



B-tree: data pointers stored in all nodes

B+ tree:

- Data pointers stored only in leaf nodes
- The leaf nodes are linked

Outline

B+ tree data structures

B+ tree operations

- Search
- Insertion
- Deletion

Primary vs. secondary indexes

B+ Tree Operations

A B+ tree supports the following operations:

- equality search
- range search
- insert
- delete
- bulk loading

Outline

B+ tree data structures

B+ tree operations

- **Search**
- Insertion
- Deletion

Primary vs. secondary indexes

Search

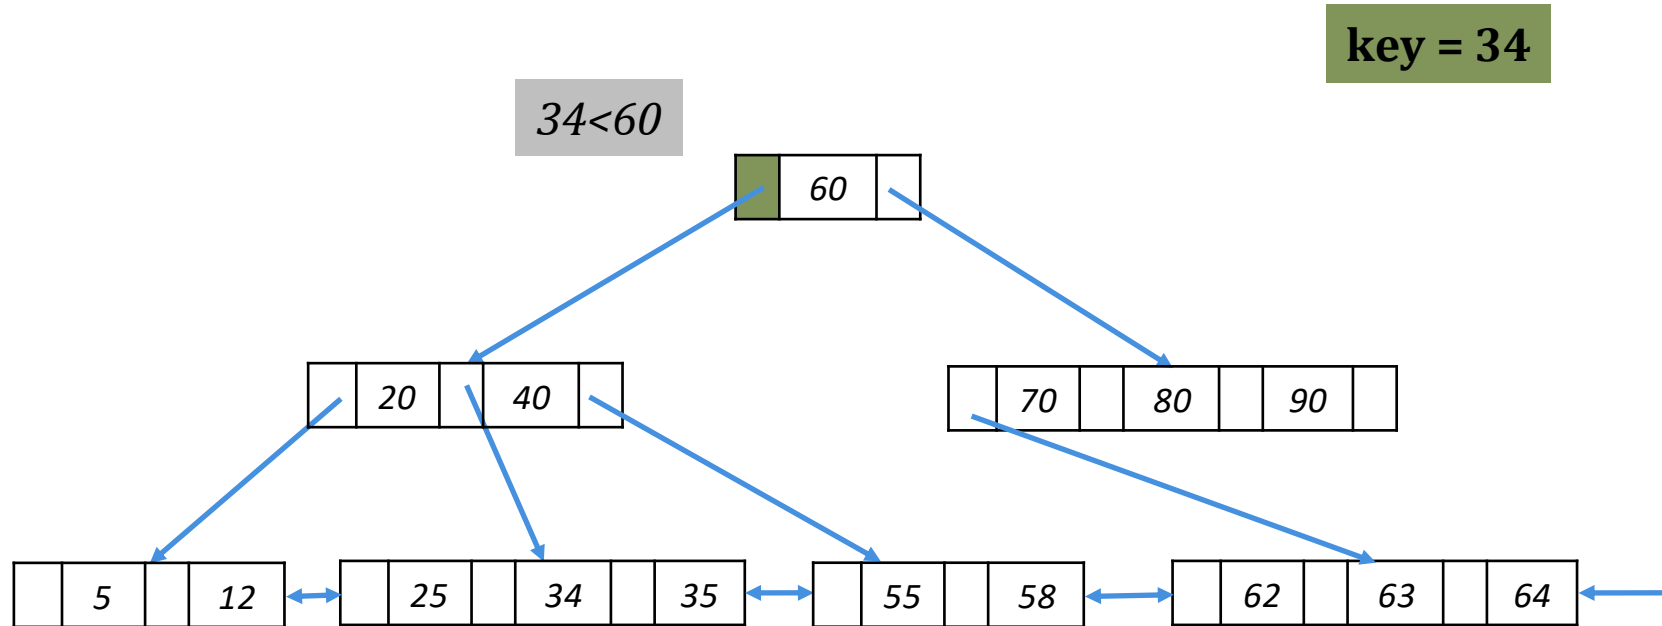
Start from the root node

Examine the index entries in non-leaf nodes to find the correct child

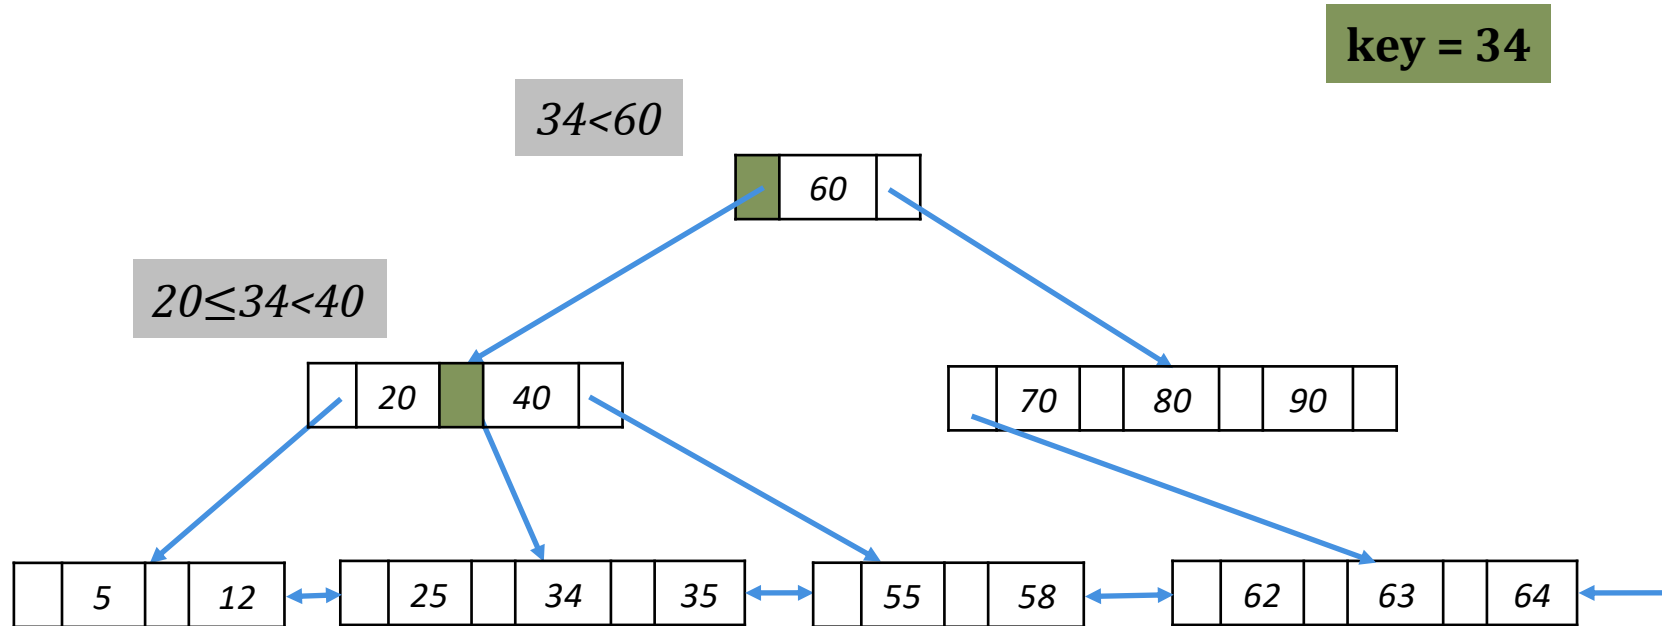
Traverse down the tree until a leaf node is reached

- For equality search, we are done
- For range search, traverse the leaves sequentially using the previous/next pointers

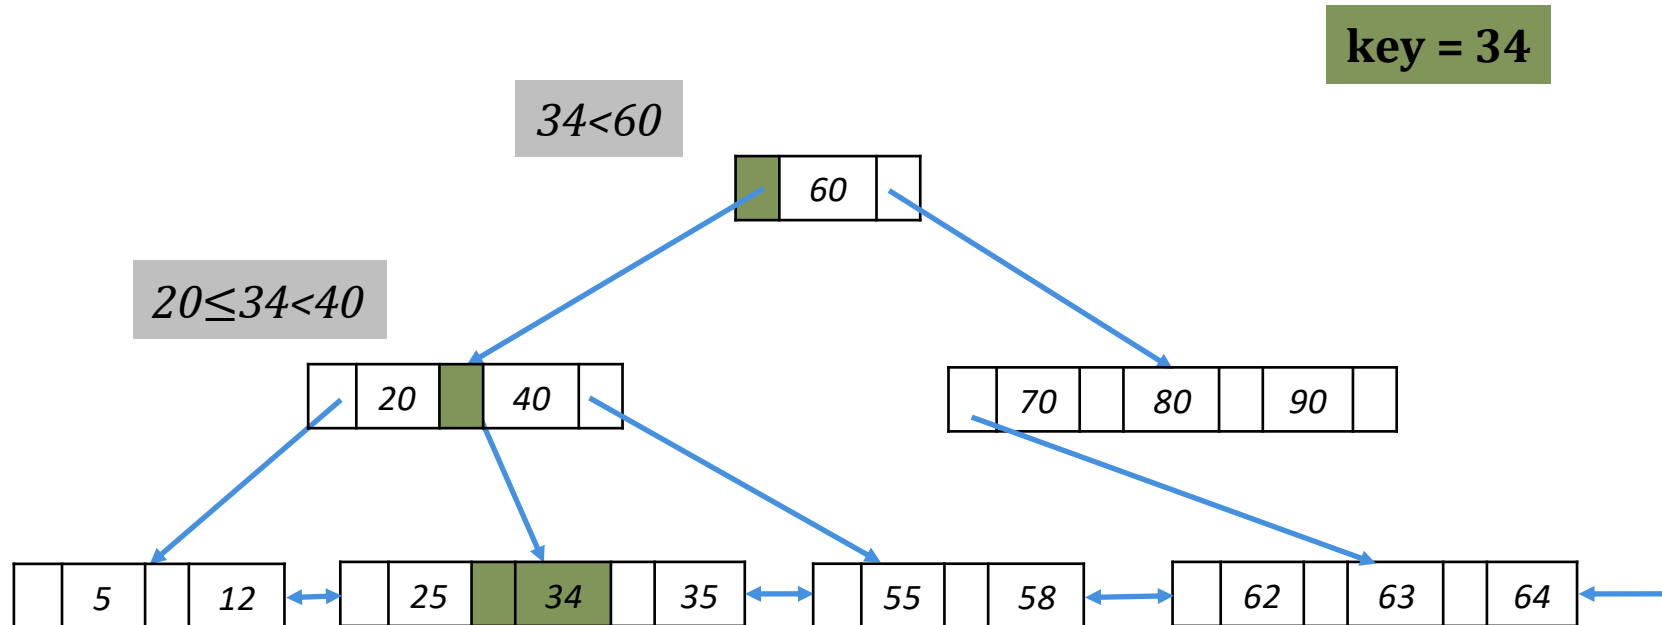
Equality Search – Example



Equality Search – Example

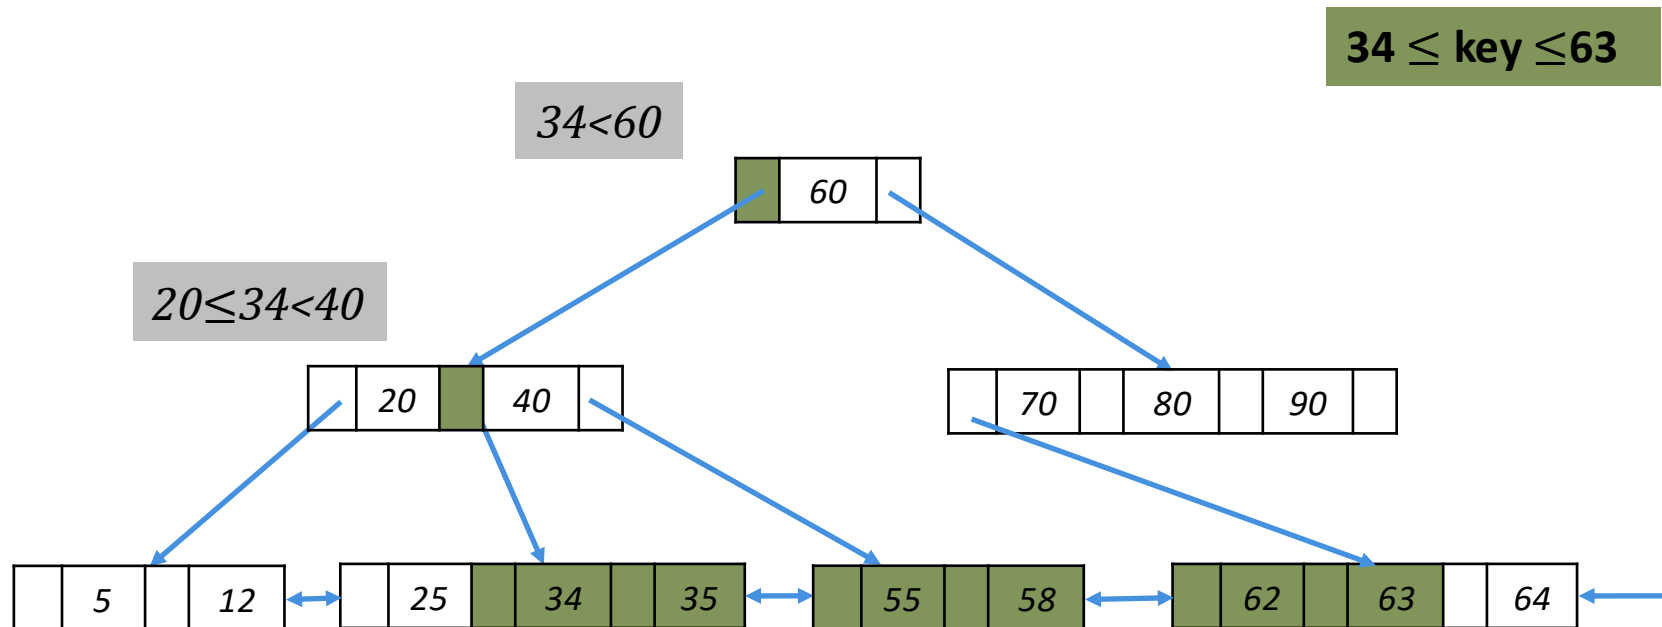


Equality Search – Example



To locate the correct data entry in the leaf node, we can do either linear or binary search

Range Search – Example



After we find the leftmost point of the range, we traverse sequentially!

Outline

B+ tree data structures

B+ tree operations

- Search
- **Insertion**
- Deletion

Primary vs. secondary indexes

Insert

Find the leaf node **L** where the entry belongs

Insert data entry in **L**

- If **L** has enough space, DONE!
- Otherwise, we must **split** **L** (into **L** and a new node **L'**)


This can propagate **recursively** to other nodes!

- To split a non-leaf node, redistribute entries evenly, but **push up** the middle key

Insert Example

order $d = 2$


Insert 15



Insert Example

order $d = 2$


Insert 15



Insert Example

order $d = 2$


Insert 8



Insert Example

order $d = 2$

Insert 8




the leaf node is full so
we must split it!

Insert Example

order $d = 2$

Insert 8

The middle key (**5**) must be inserted to the parent node




Insert Example

order $d = 2$

Insert 8

The middle key (**5**) must be inserted to the parent node




Insert Example – Parent Node Split

order $d = 2$

What if the parent node is already full before the insertion?

Insert 8




Insert Example – Parent Node Split

order $d = 2$

What if the parent node is already full before the insertion?

Insert 8

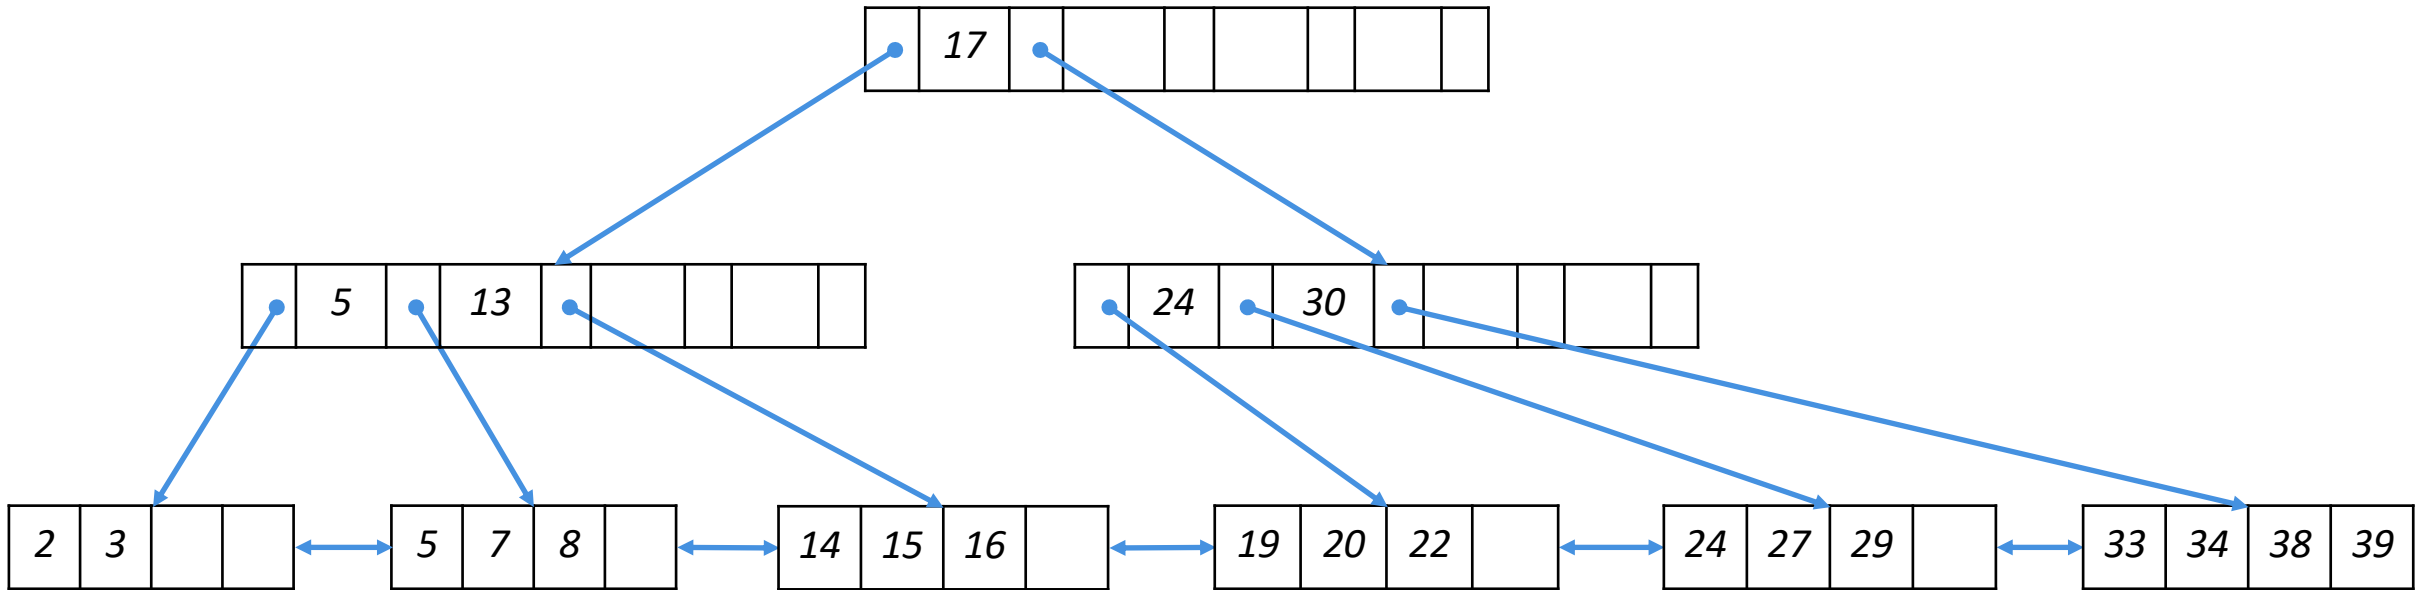
The parent node must also be split



Insert Example – Parent Node Split

order $d = 2$

Insert 8



Insert Properties

The B+ tree insertion algorithm has several attractive qualities:

- About the same cost as equality search
- ***Self-balancing***: the tree remains balanced (with respect to height) even after multiple insertions

Outline

B+ tree data structures

B+ tree operations

- Search
- Insertion
- **Deletion**

Primary vs. secondary indexes

Delete

Find the leaf node **L** where the entry belongs

Remove the entry

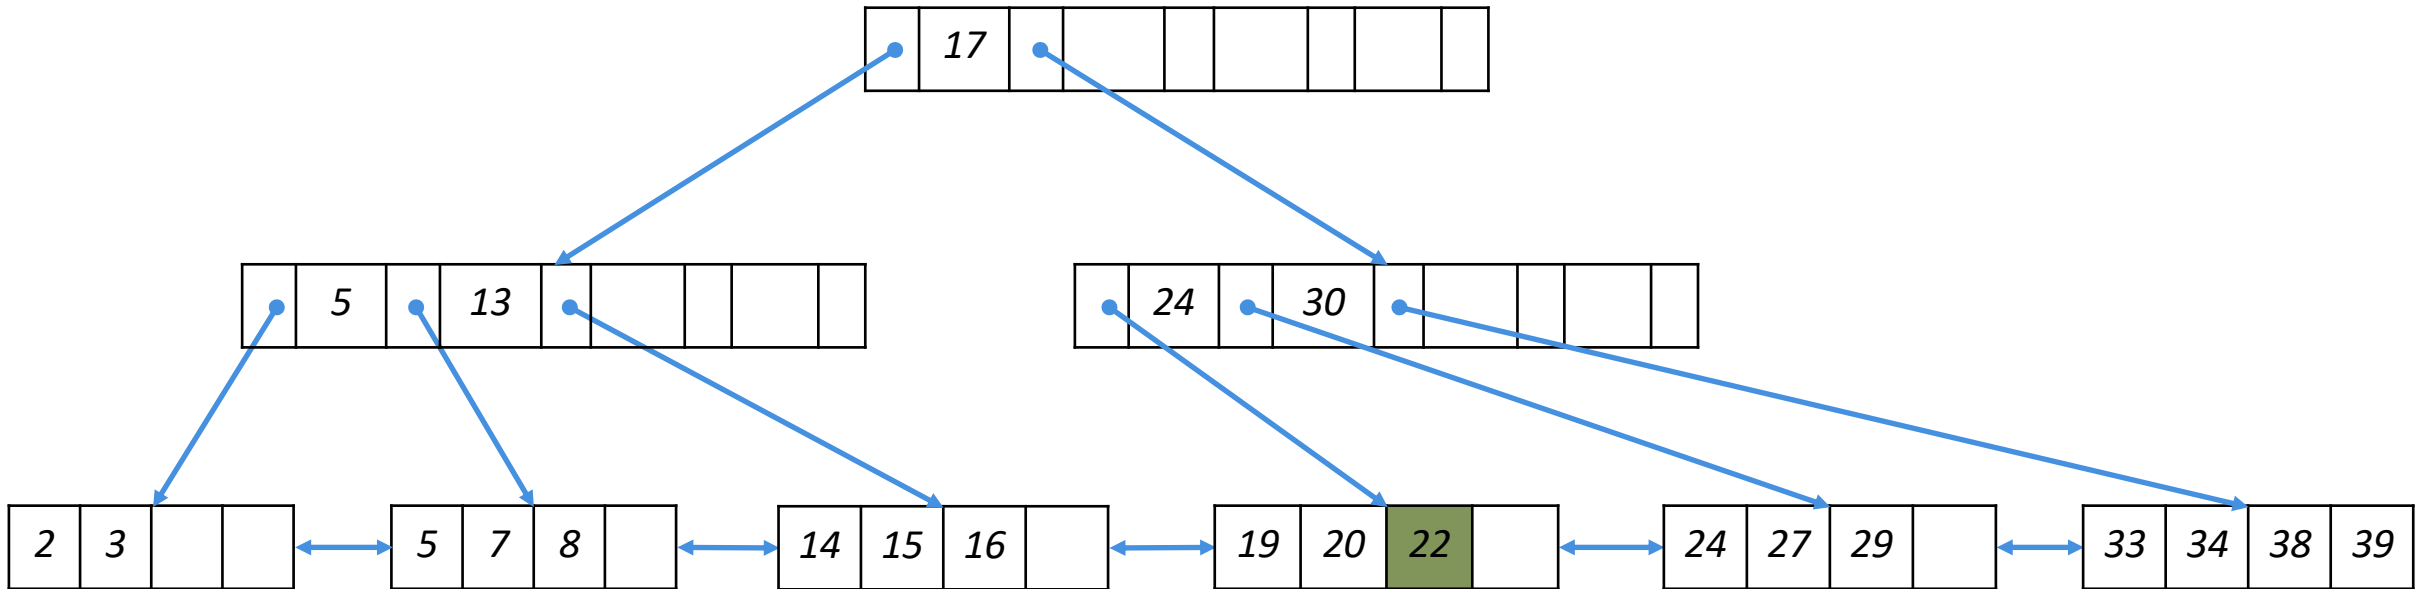
- If **L** is at least half-full, DONE!
- If **L** has only $d-1$ entries,
 - Try to **redistribute** borrowing entries from a neighboring **sibling**
 - If redistribution fails, **merge L** and sibling

If a merge occurred, we must delete an entry from the parent of **L**

Delete – Example

order $d = 2$

Delete 22

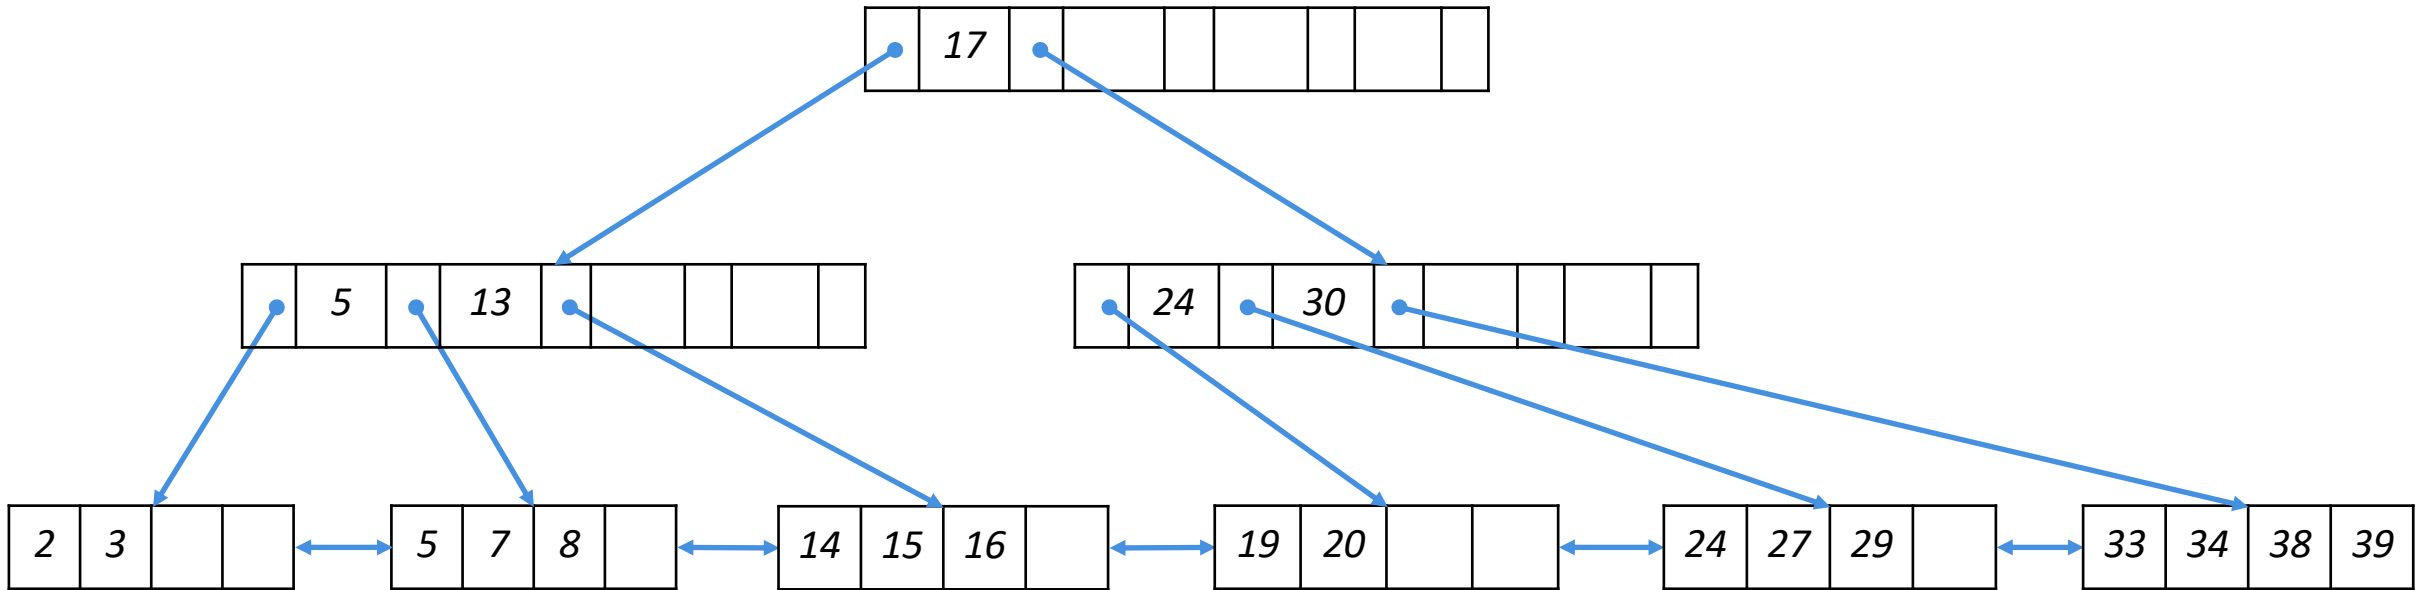


Since by deleting 22 the node remains half-full, we simply remove it

Delete – Example

order $d = 2$

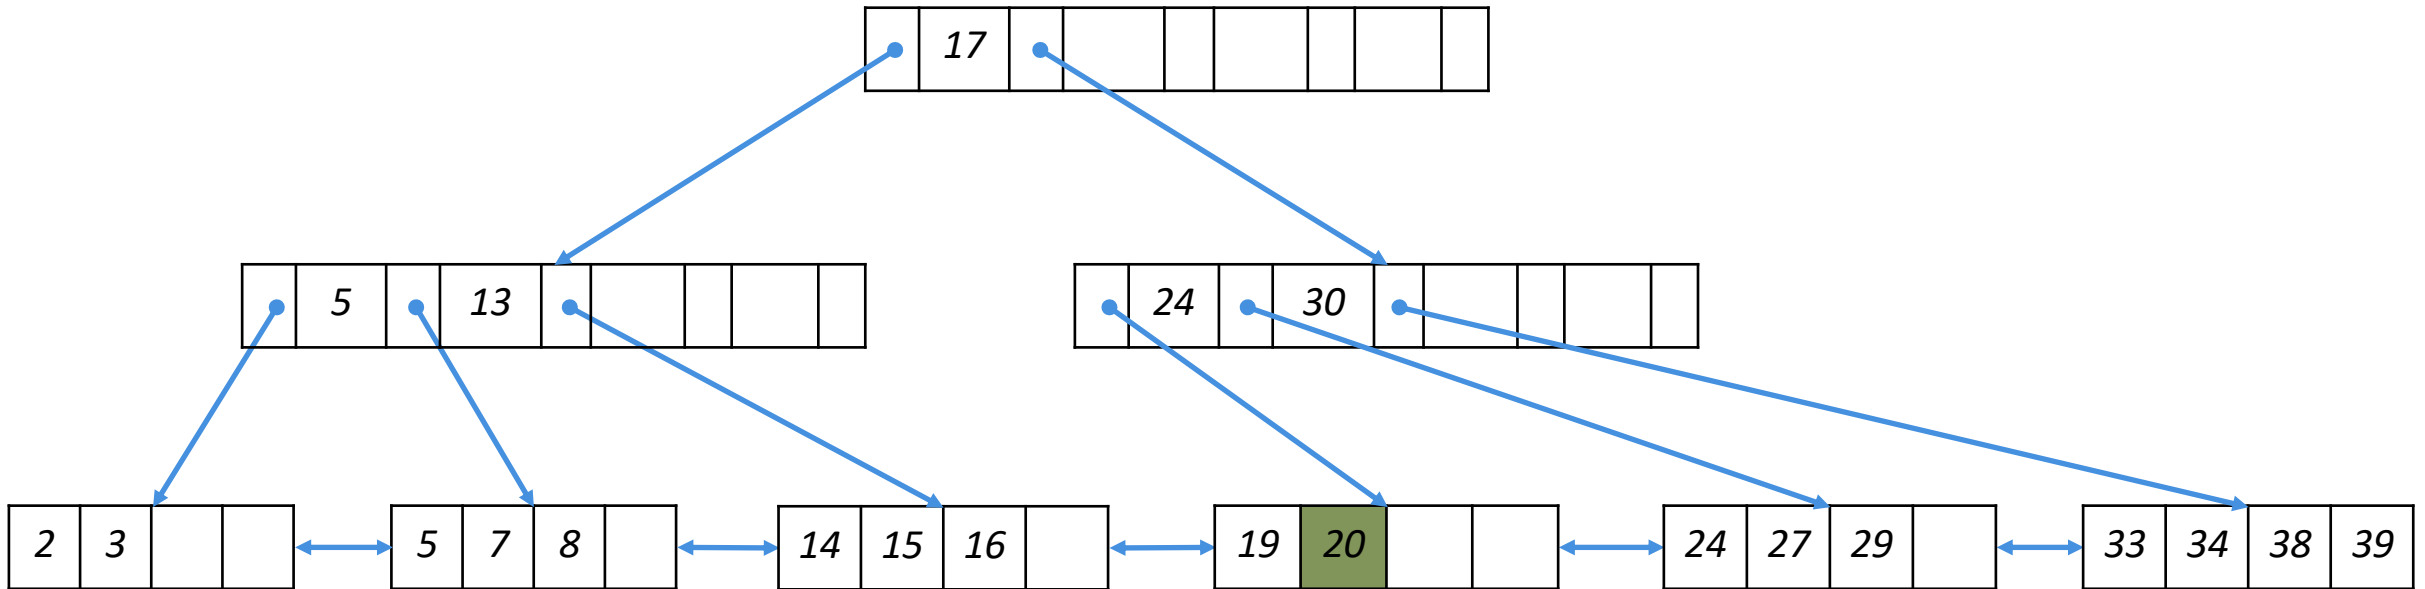
Delete 22



Delete – Example

order $d = 2$

Delete 20

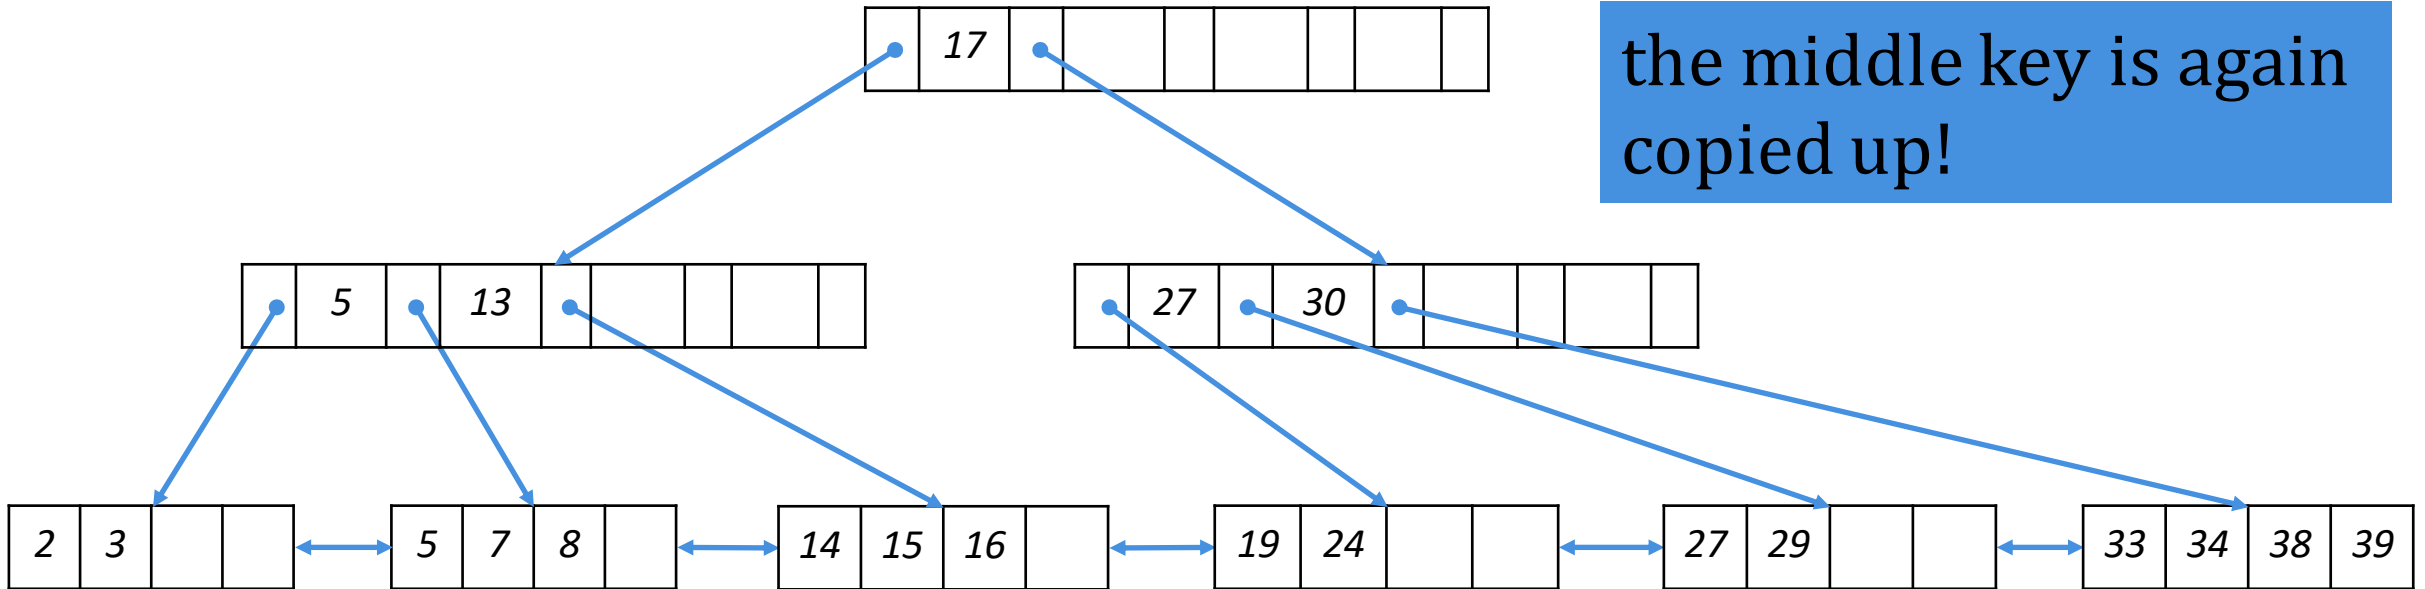


by removing 20 the node is not half-full anymore, so we attempt to redistribute!

Delete – Example

order $d = 2$

Delete 20



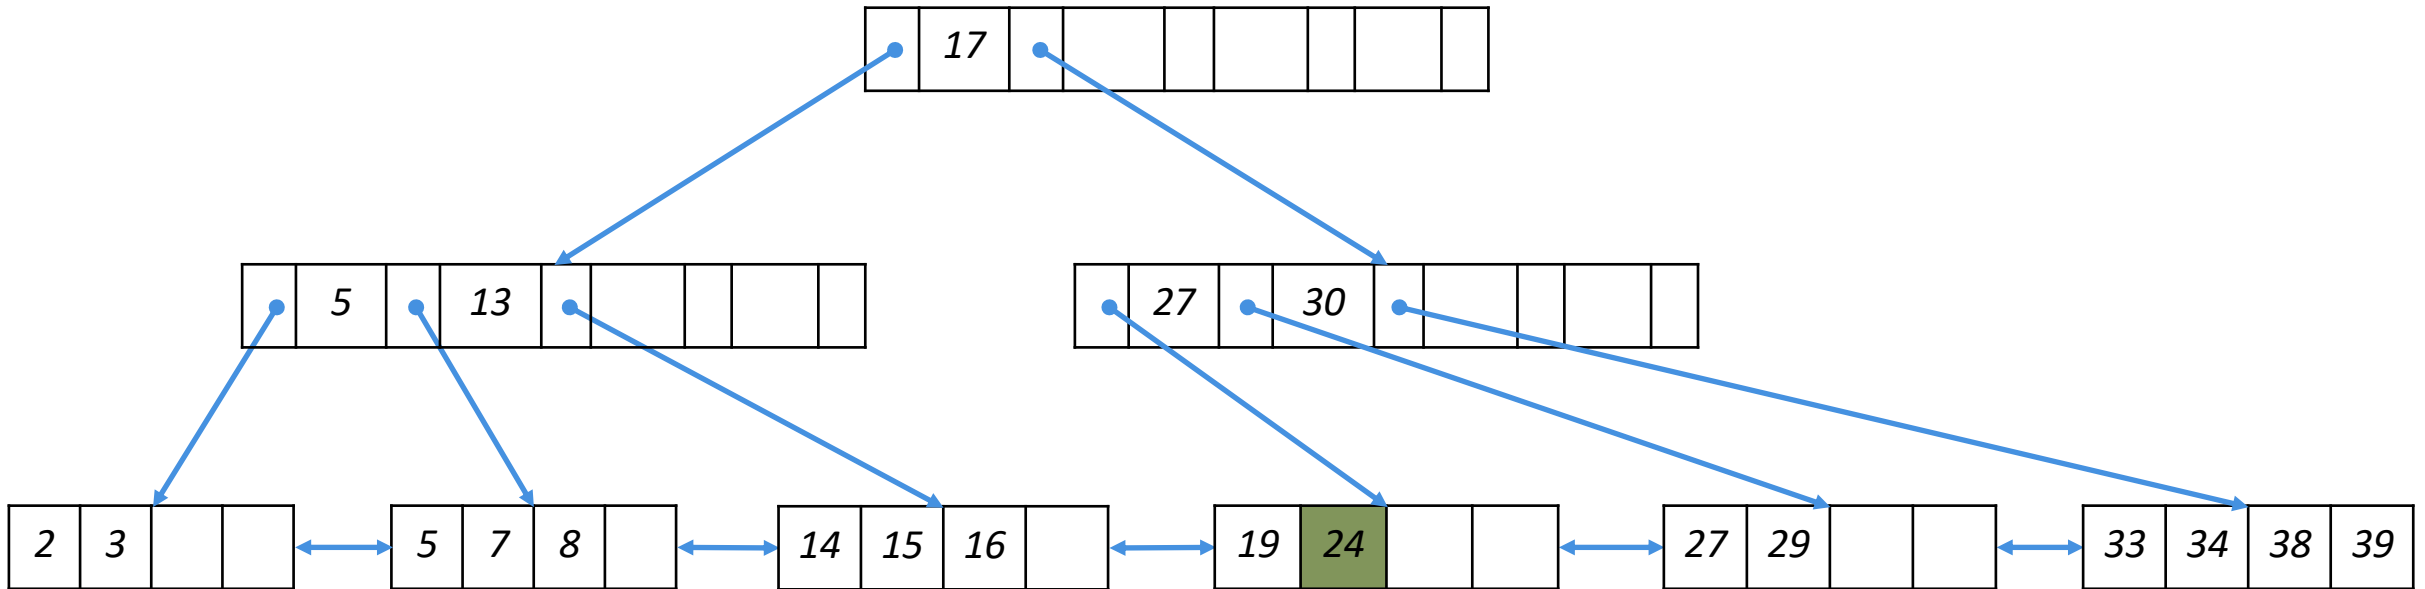
the middle key is again copied up!

by removing 20 the node is not half-full anymore, so we attempt to redistribute!

Delete – Example

order $d = 2$

Delete 24

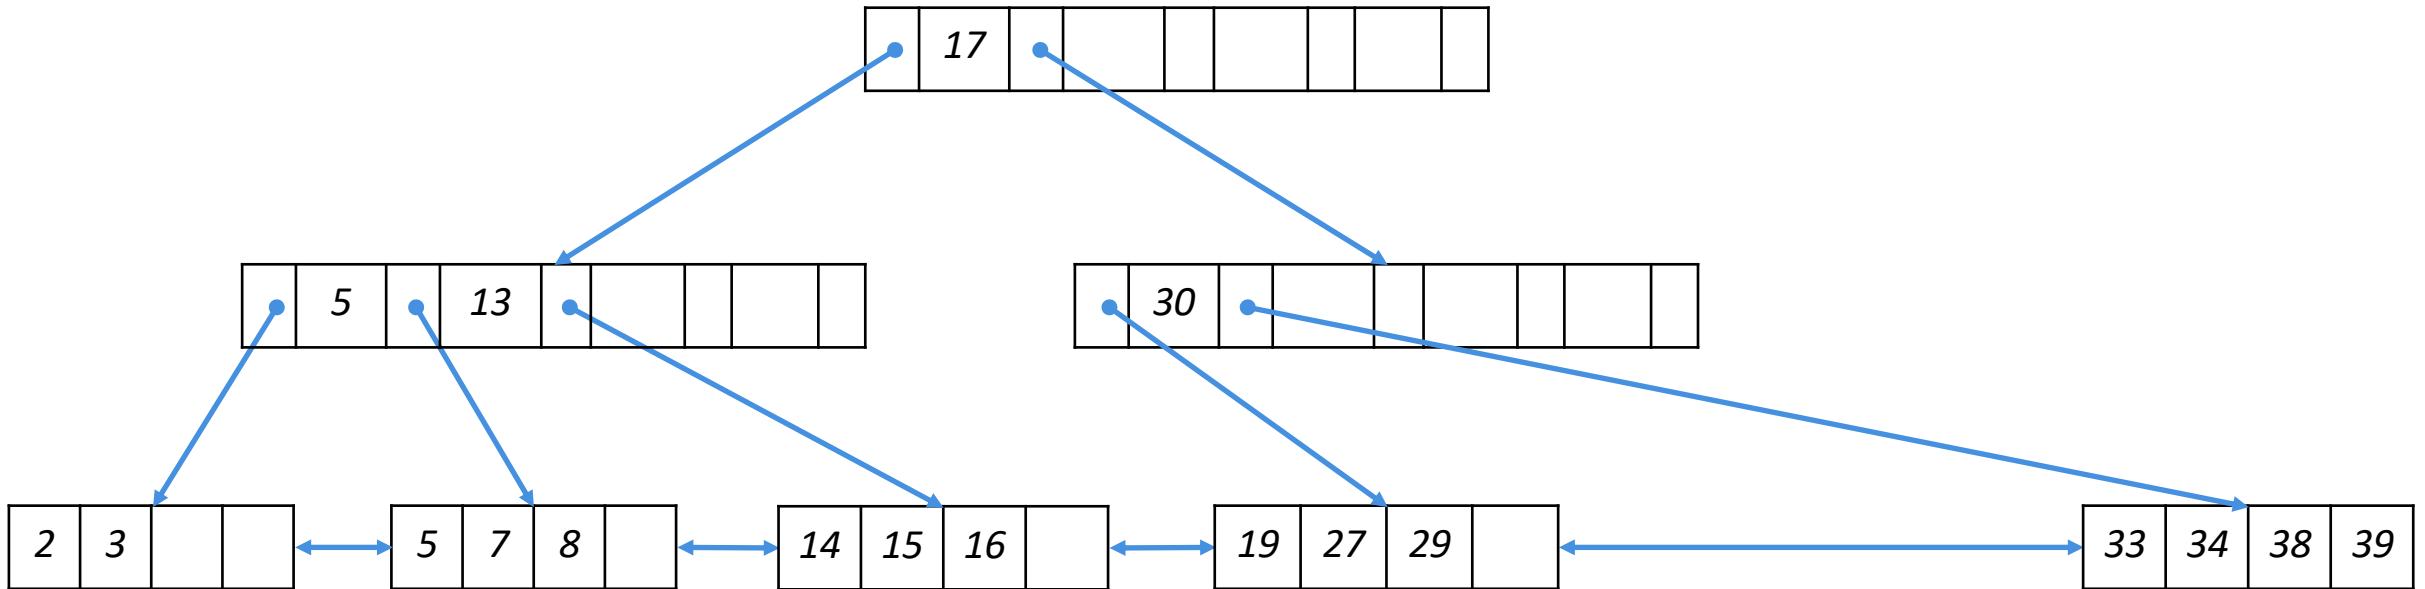


in this case, we have to merge nodes!

Delete – Example

order $d = 2$

Delete 24




in this case, we have to merge nodes!

Delete – Example

order $d = 2$

Delete 24




We are not done, since the resulting non-leaf node is not half-full!

Delete – Example

order $d = 2$

Delete 24

Merge non-leaf nodes



We are not done, since the resulting non-leaf node is not half-full!

More on Delete

Redistribution of entries is also possible for the non-leaf nodes

We can also try to redistribute using *all siblings*, and not only the neighboring one

In real systems, deletion is often implemented without adjusting the tree structure, because files typically grow rather than shrink

Outline

B+ tree data structures

B+ tree operations

- Search
- Insertion
- Deletion

Primary vs. secondary indexes

Primary vs. Secondary Index

If the search key contains the primary key, it is called a **primary index**

- In a primary index, there are no duplicates for a value of the search key
- There can only be one primary index!


Any other index is called a **secondary index**

If the search key contains a candidate key, it is called a **unique index**

- A primary index is also a unique index

Alternative Secondary Index Design

Primary Index




Secondary Index


Primary index stores **records**, secondary index stores **rid**

Alternative Secondary Index Design

Primary Index



Secondary Index




Primary index stores **records**, secondary index stores **rid**



Both primary and secondary indexes store **rid**

Alternative Secondary Index Design

Primary Index



Secondary Index



Primary index stores **records**, secondary index stores **rid**

Both primary and secondary indexes store **rid**

Secondary index store the **primary key**

- Only primary index is changed when updating rid

Summary

B+ tree data structures

B+ tree operations

- Search
- Insertion
- Deletion

Primary vs. secondary indexes