



CS 564: Database Management Systems

Lecture 27: Query Optimization I

Xiangyao Yu
4/1/2024

Announcement

Midterm grade announced

Regrade request should be submitted (to Gradescope) **by next Monday (April 8th)**

Module B3 Query Processing

Relational operators I

Relational operators II

Query optimization I

Query optimization II

Column Store

Outline

Optimization example

- Pipelining
- Selection pushdown
- Projection pushdown

Plan enumeration

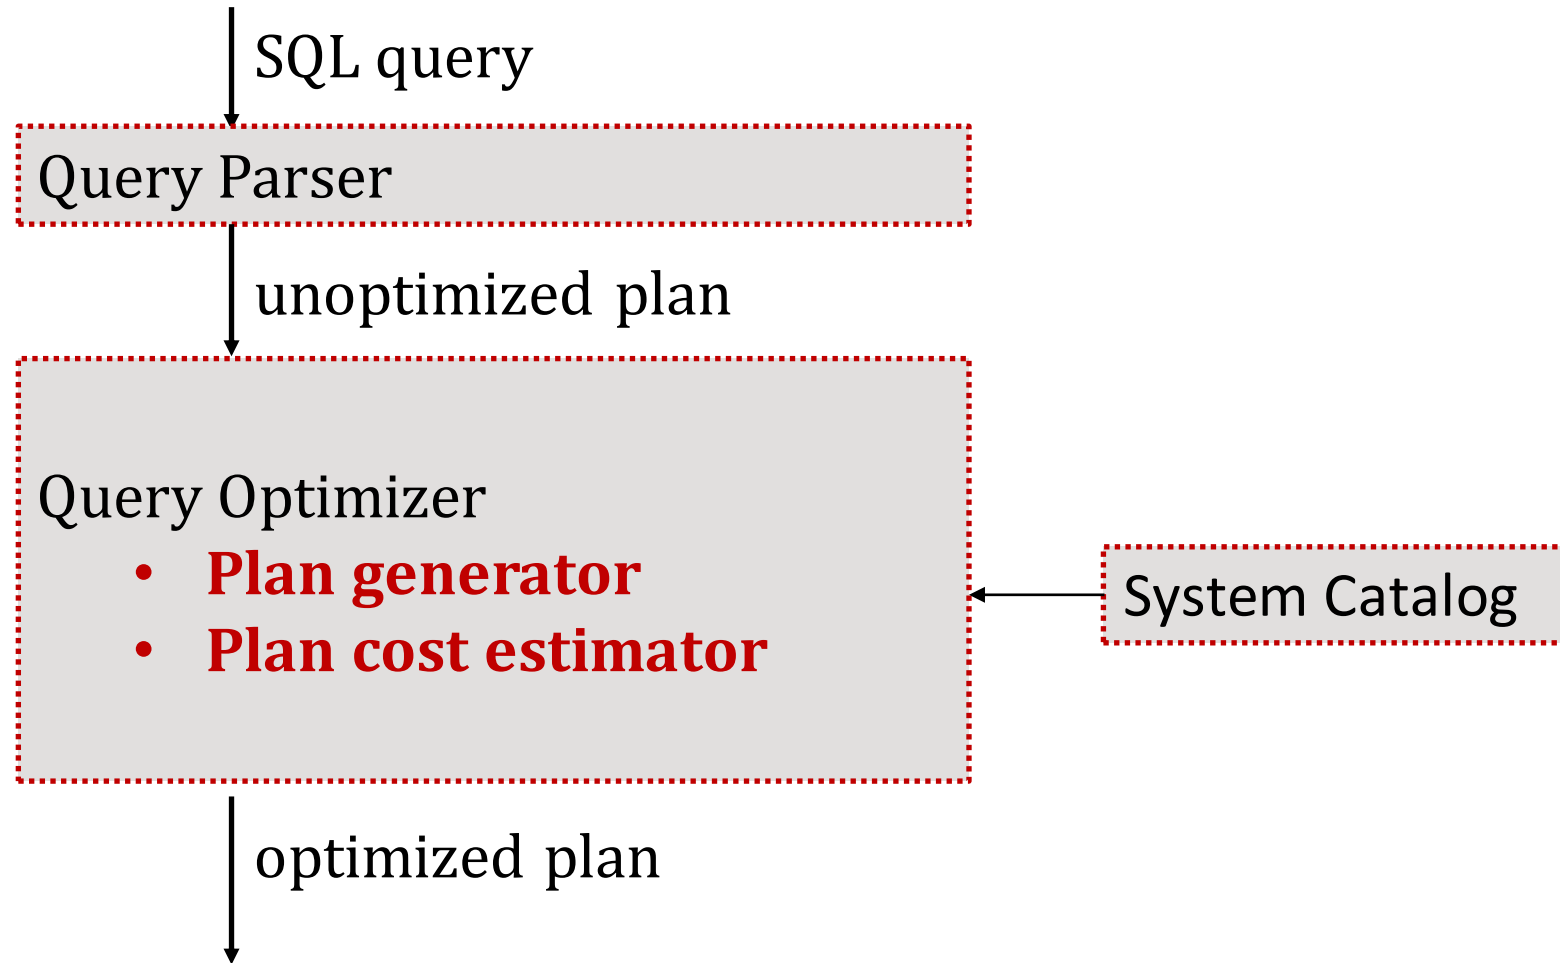
Same Query, Different Cost

Relational query language allows a database to evaluate a query in many different ways


Difference in cost between best and worst plans may be several orders of magnitude

Query optimizer: generates alternative plans and choose the plan with the least estimated cost

Architecture of an Optimizer



Example



```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```

	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

Convention: **left child is the outer table**

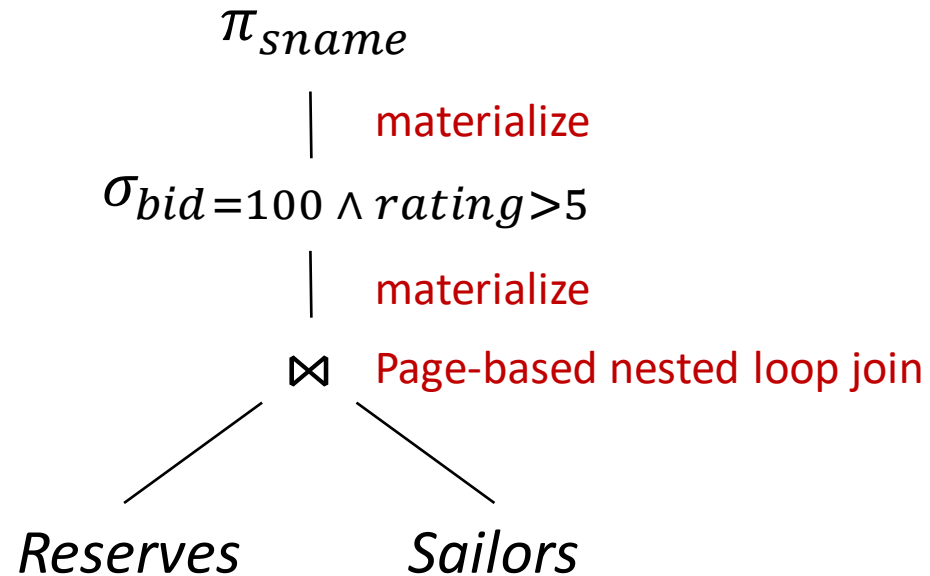
Assume 1% tuples in Reserves satisfy $bid = 100$ and 50% tuples in Sailors satisfy $rating > 5$

Example – Plan 1

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```

	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5




Materialize: create a temporary table to hold the intermediate results

Example – Plan 1

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```

	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid =100
- 50% tuples in Sailors satisfy rating > 5



Join IO cost:

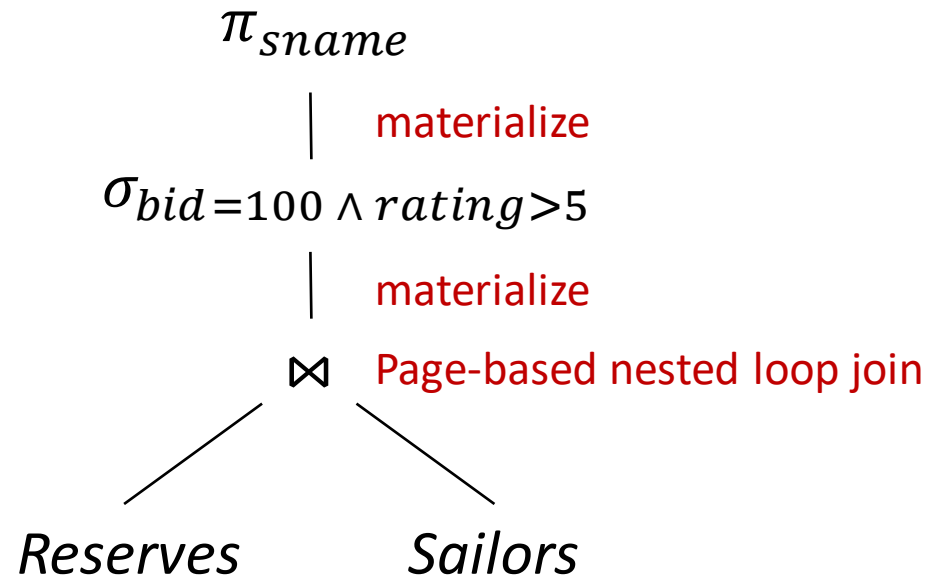
$$M_R + M_R \times M_S = 1000 + 1000 \times 500 = 501,000$$

Example – Plan 1

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```

	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5



Join IO cost:

$$M_R + M_R \times M_S = 1000 + 1000 \times 500 = 501,000$$

Join output size:

$$100 \times 1000 \times (40 B + 50 B) / 4kB = 2250 \text{ pages}$$


Each tuple in Reserves joins with one tuple in Sailors (primary key and foreign key)

Example – Plan 1

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```

	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5



Join IO cost:

$$M_R + M_R \times M_S = 1000 + 1000 \times 500 = 501,000$$

Join output size:

$$100 \times 1000 \times (40 B + 50 B) / 4kB = 2250 \text{ pages}$$

Selection output size:


$$2250 \times 0.01 \times \frac{1}{2} = \sim 11 \text{ pages}$$

Example – Plan 1

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```

	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5



Join IO cost:

$$M_R + M_R \times M_S = 1000 + 1000 \times 500 = 501,000$$

Join output size:

$$100 \times 1000 \times (40 B + 50 B) / 4kB = 2250 \text{ pages}$$

Selection output size:

$$2250 \times 0.01 \times \frac{1}{2} = \sim 11 \text{ pages}$$

Total IO cost:

$$501,000 + 2250 \times 2 + 11 \times 2 = \mathbf{505,522 \text{ IOs}}$$

Outline

Optimization example


- **Pipelining**
- Selection pushdown
- Projection pushdown

Plan enumeration


Pipelined Evaluation

Result of one operator pipelined to another operator tuple-by-tuple without materializing the intermediate relation

Examples:



Results of first join pipelined
into the second join

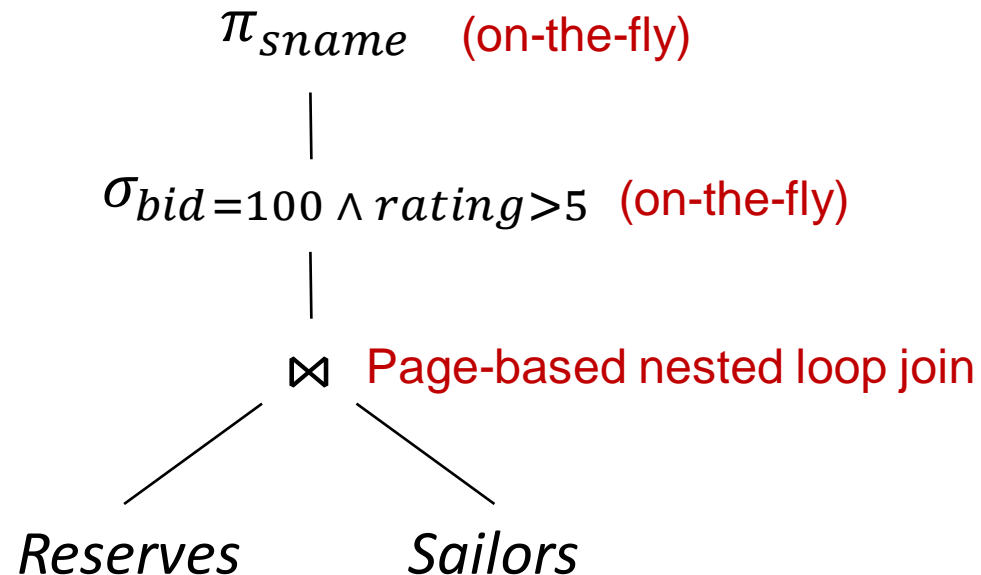


Example – Plan 2

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```

	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5




Join IO cost:

$$M_R + M_R \times M_S = 1000 + 1000 \times 500 = 501,000$$

Total IO cost = Join cost = **501,000** IOs

Example – Plan 3 (Pipelined Joins)



$R \bowtie S$ IO cost:

$$50 + 100 \times 50 = 5050 \text{ IOs}$$

Write intermediate results

$$150 \text{ IOs}$$

IO cost to join *T*:

$$150 + 150 \times 30 = 4650 \text{ IOs}$$

Total: 9850 IOs

Pipelined join does not write intermediate results back to disk and saves **150×2 IOs**

Outline

Optimization example

- Pipelining
- **Selection pushdown**
- Projection pushdown

Plan enumeration

Selection Pushdown


A join is relatively expensive

Good heuristic: **reduce join table sizes** as much as possible

One approach: apply selection early

Example – Plan 4

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```




	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid =100
- 50% tuples in Sailors satisfy rating > 5

Example – Plan 4

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```



	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500


- 1% tuples in Reserves satisfy bid =100
- 50% tuples in Sailors satisfy rating > 5

Selection on Reserves

– input: 1000 pages, output: 10 pages

Example – Plan 4

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```



	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5

Selection on Reserves


- input: 1000 pages, output: 10 pages

Selection on Sailors

- input: 500 pages, output: 250 pages

Example – Plan 4

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```



	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5

Selection on Reserves

- input: 1000 pages, output: 10 pages

Selection on Sailors


- input: 500 pages, output: 250 pages

Nested-loop join

- $10 + 10 \times 250 = 2510$ IOs

Example – Plan 4

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```



	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5

Selection on Reserves

- input: 1000 pages, output: 10 pages

Selection on Sailors

- input: 500 pages, output: 250 pages

Nested-loop join

- $10 + 10 \times 250 = 2510$ IOs

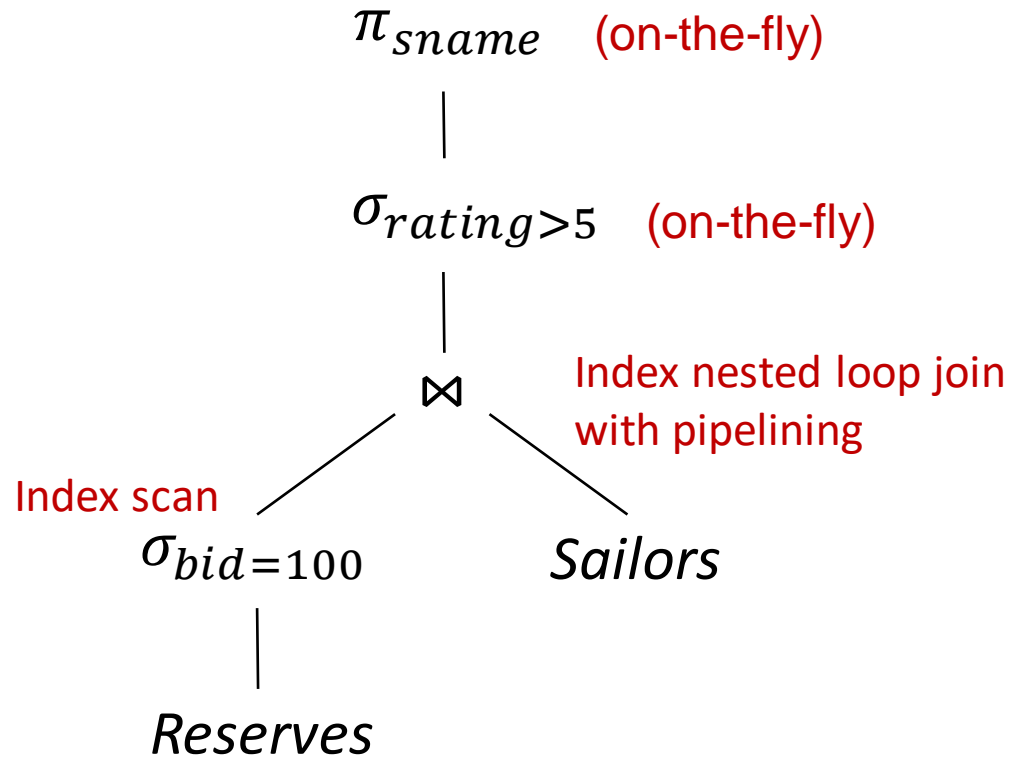
Total IO cost

= selection cost + join cost

= $(1000+10+500+250) + (2510) = \mathbf{4270}$ IOs

Example – Plan 5 (Using Index)

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```




	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy $bid = 100$
- 50% tuples in Sailors satisfy $rating > 5$
- **Clustered hash index on R.bid and S.sid**

Example – Plan 5 (Using Index)

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```



	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5
- **Clustered hash index on R.bid and S.sid**

Index scan on Reserves

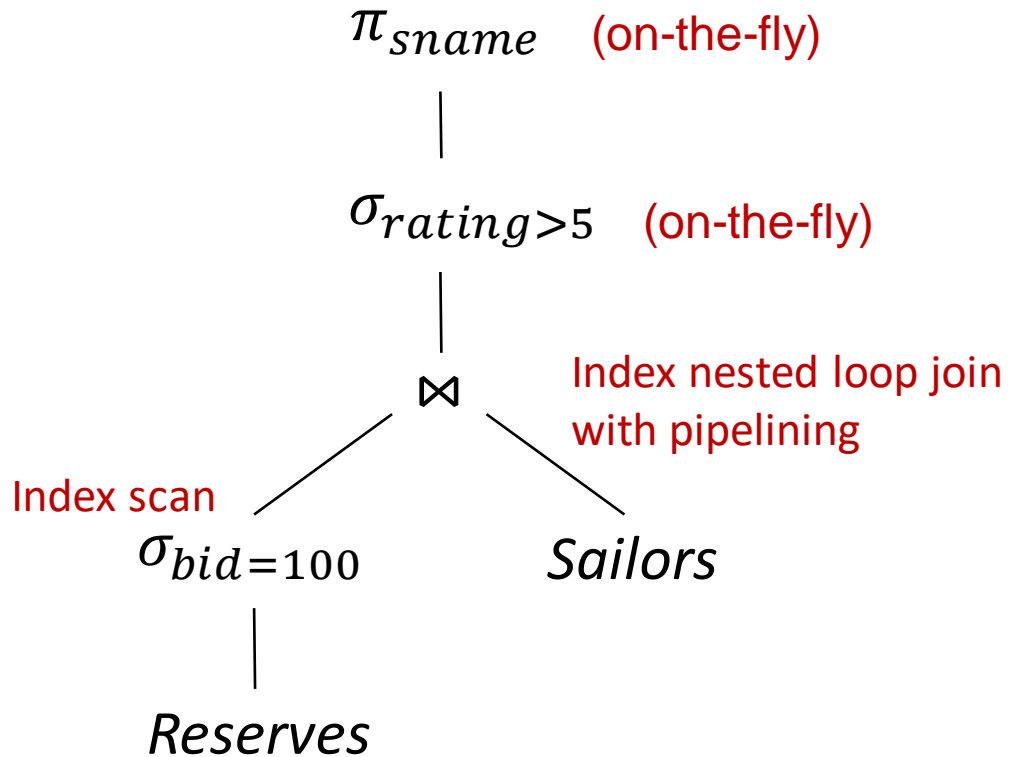
– 10 pages

Index lookup on Sailors

– 1000 IOs (assume 1 IO per lookup)

Example – Plan 5 (Using Index)

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 100
      AND S.rating > 5;
```



	Tuple size	#tuples per page	# of pages
Reserves	40 bytes	100	1000
Sailors	50 bytes	80	500

- 1% tuples in Reserves satisfy bid = 100
- 50% tuples in Sailors satisfy rating > 5
- **Clustered hash index on R.bid and S.sid**

Index scan on Reserves

– 10 pages

Index lookup on Sailors

– 1000 IOs (assume 1 IO per lookup)

Total IO cost = **1100** IOs

Selection pushdown is not always a good idea

Outline

Optimization example

- Pipelining
- Selection pushdown
- **Projection pushdown**

Plan enumeration

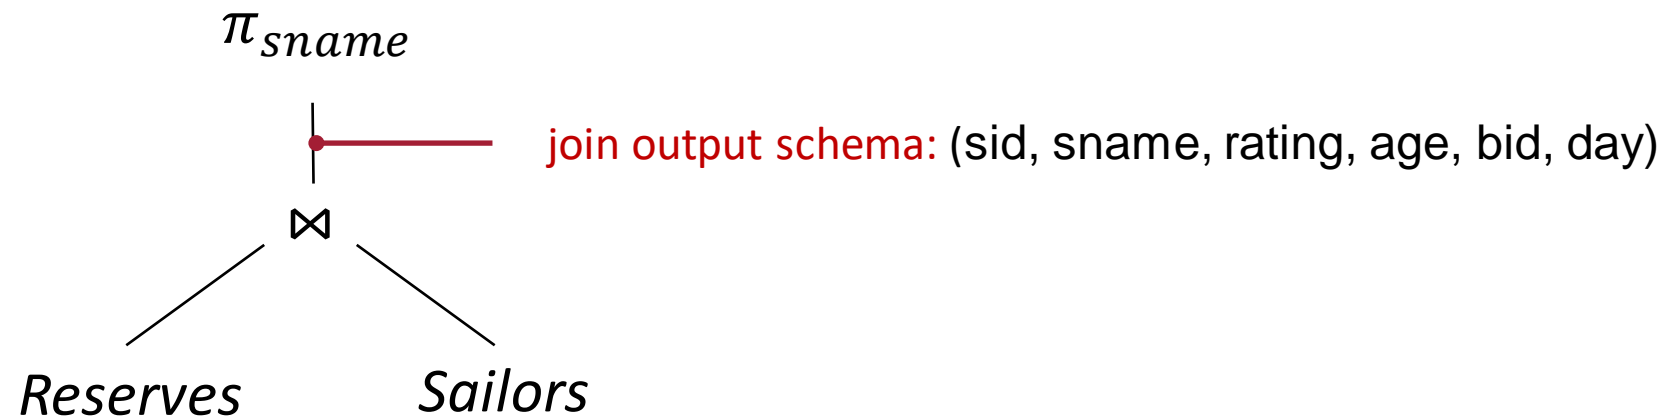
Projection Pushdown

Apply projection early to reduce table size

Example:

Sailors (sid, sname, rating, age)

Reserves (sid, bid, day)





Projection Pushdown

Apply projection early to reduce table size

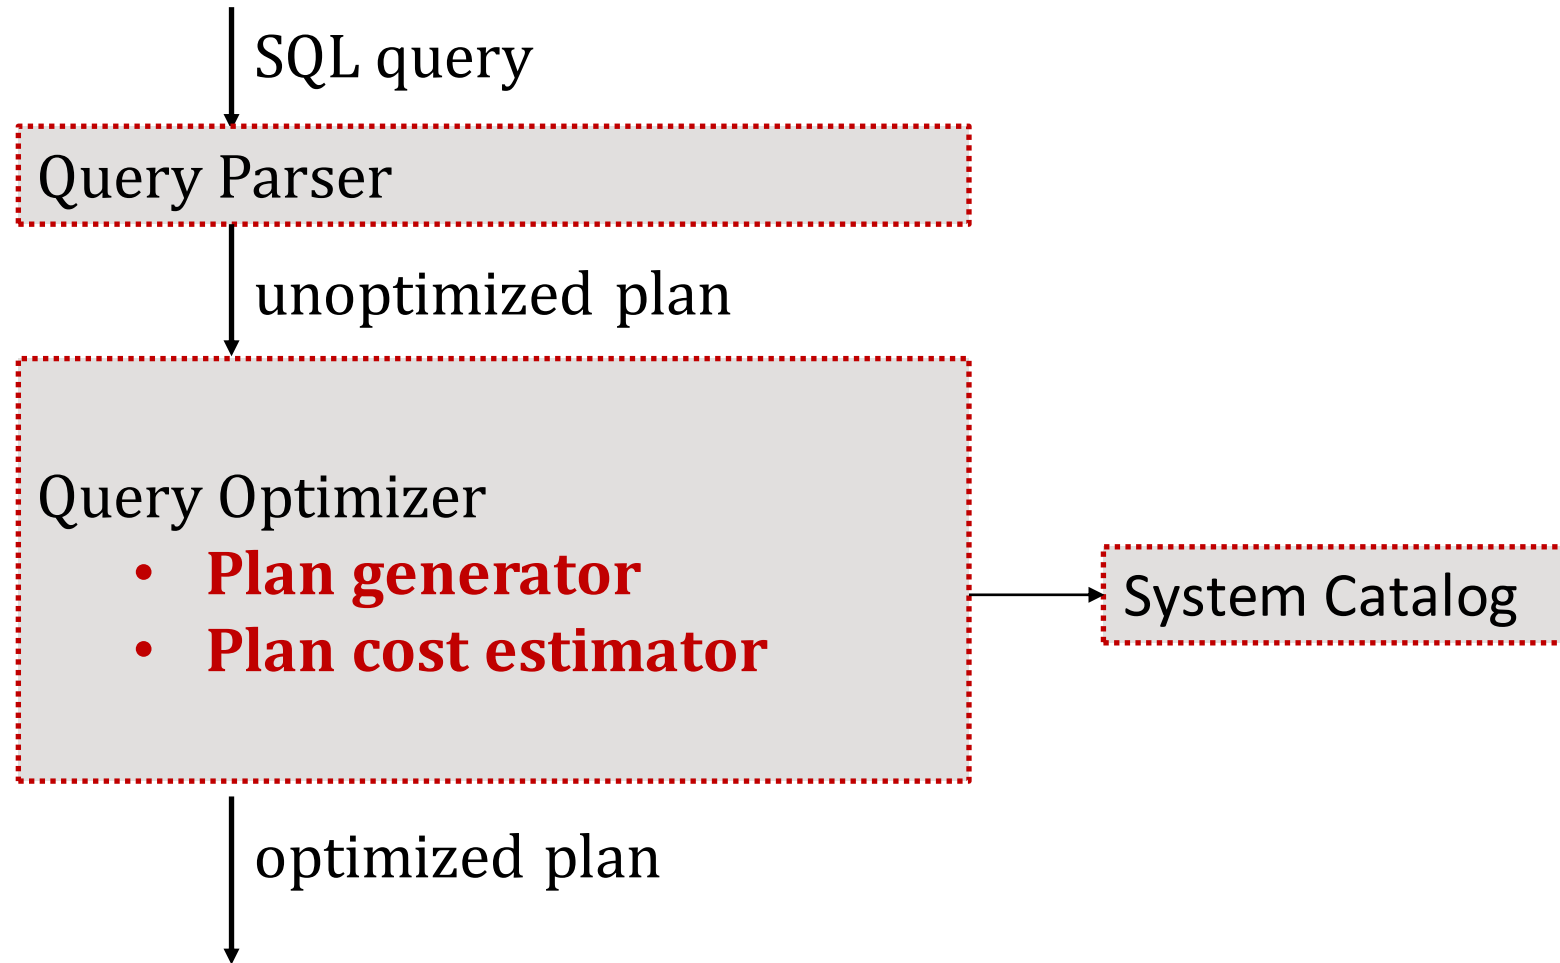
Example:

Sailors (sid, sname, rating, age)

Reserves (sid, bid, day)



Query Optimizer



Outline

Optimization example

- Pipelining
- Selection pushdown
- Projection pushdown

Plan enumeration

Plan Generation

The space of possible query plans is huge and it is hard to navigate

- Access path (file scan, index scan)
- Join algorithm (merge-sort, hash, nested-loop, index join)
- Join order
- Push down selection & projection
- Sorting
- etc.

Relational Algebra provides us with rules that transform one RA expression to an equivalent one

- Push down selections & projections
- Join reordering

These transformations allow us to construct many alternative query plans

Equivalences – Selection

Cascading of selections

$$\sigma_{c_1 \wedge c_2 \wedge \dots \wedge c_n} \equiv \sigma_{c_1}(\sigma_{c_2}(\dots(\sigma_{c_n}(R))\dots))$$

Selections are commutative

$$\sigma_{c_1}(\sigma_{c_2}(R)) = \sigma_{c_2}(\sigma_{c_1}(R))$$

Equivalences – Projection

Cascading of projections

- Successively eliminating columns from a relation is equivalent to simply eliminating all but the columns retained by the final projection

$$\pi_{a_1} \equiv \pi_{a_1}(\pi_{a_2}(\dots(\pi_{a_n}(R))\dots))$$

Equivalences – Join

Joins are commutative

$$R \bowtie S \equiv S \bowtie R$$


Joins are associative

$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$$


Example

$$R \bowtie (S \bowtie T) \equiv (T \bowtie R) \bowtie S$$

Equivalences – Example 1




Equivalences – Example 2




Join Orders


$A \bowtie B \bowtie C \bowtie D$



Left-deep



Linear




Bushy

Join Orders

Many databases consider only **left-deep plans**, for the following two reasons

- As the number of joins increases, the number of alternative plans increases rapidly and it becomes necessary to prune the space of alternative plans
- Left deep trees allow us to generate **fully pipelined** plans, in which all joins are evaluated using pipelining.



Summary

Optimization example

- Pipelining
- Selection pushdown
- Projection pushdown

Plan enumeration