



CS 564: Database Management Systems

Lecture 28: Query Optimization II

Xiangyao Yu
4/3/2024

Module B3 Query Processing

Relational operators I

Relational operators II


Query optimization I

Query optimization II

Column Store

Clarification on Pipelined Hash Join

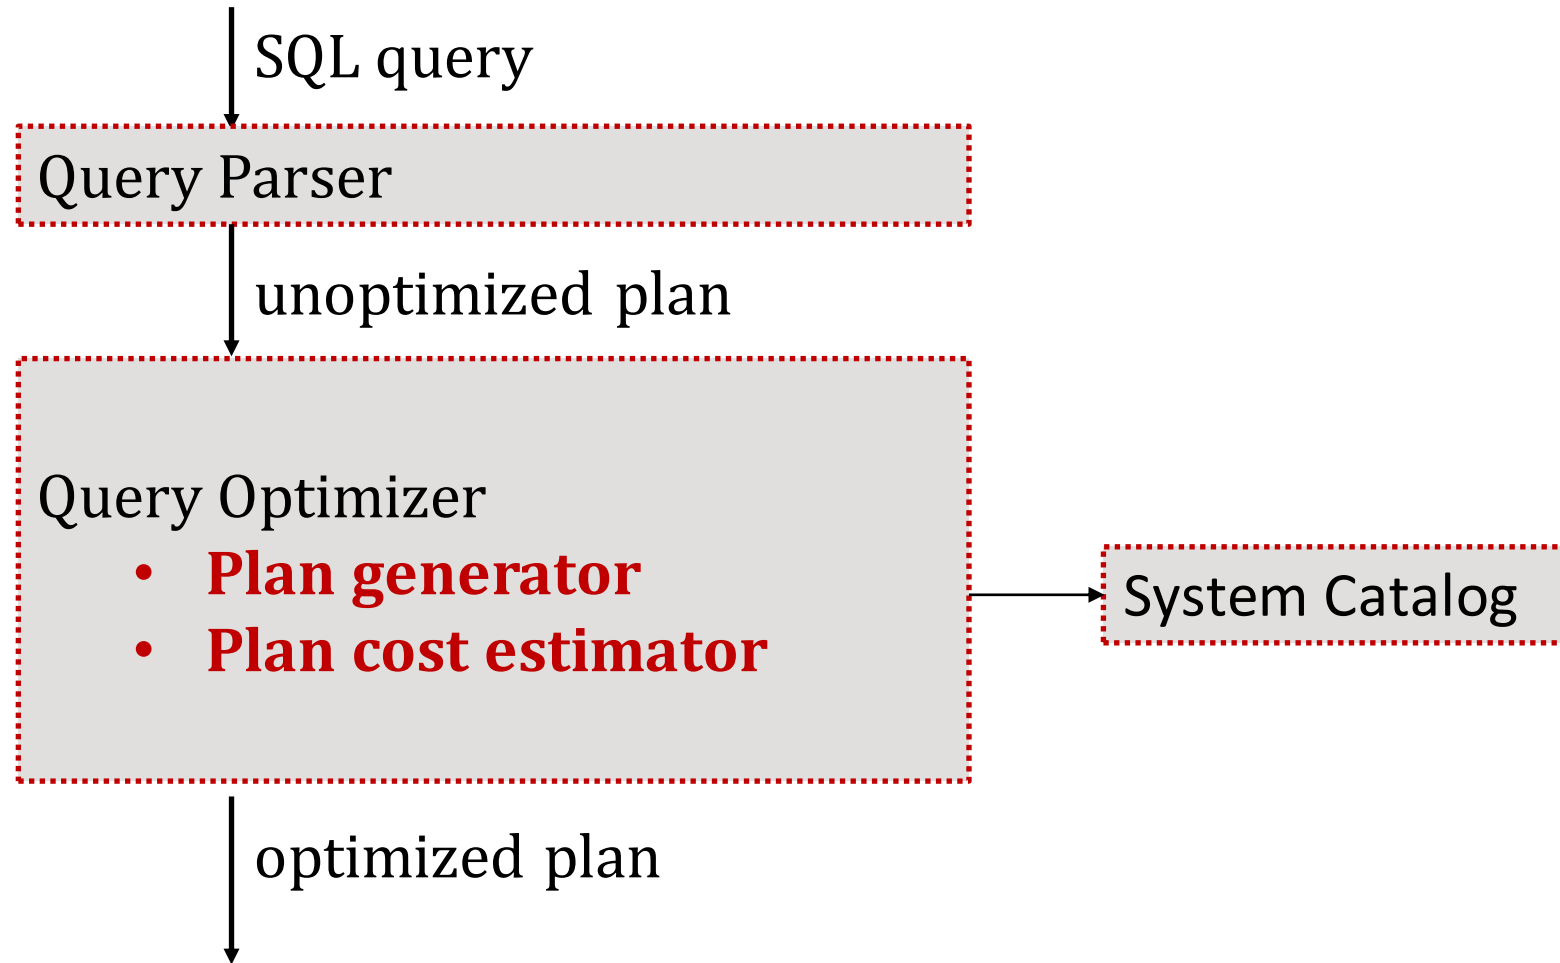
Results of one join
pipelined into the
next join



Step 1: Build hash tables on
relations B, C, and D

Step 2: Read table A tuple by tuple
and lookup hash tables in order

Architecture of an Optimizer



Outline

Aggregation pushdown

System catalog

Cost estimation

- Cardinality estimation
- IO cost estimation

Plan enumeration

Group-By Aggregation Pushdown


Partial group by can also reduce cost

Example:

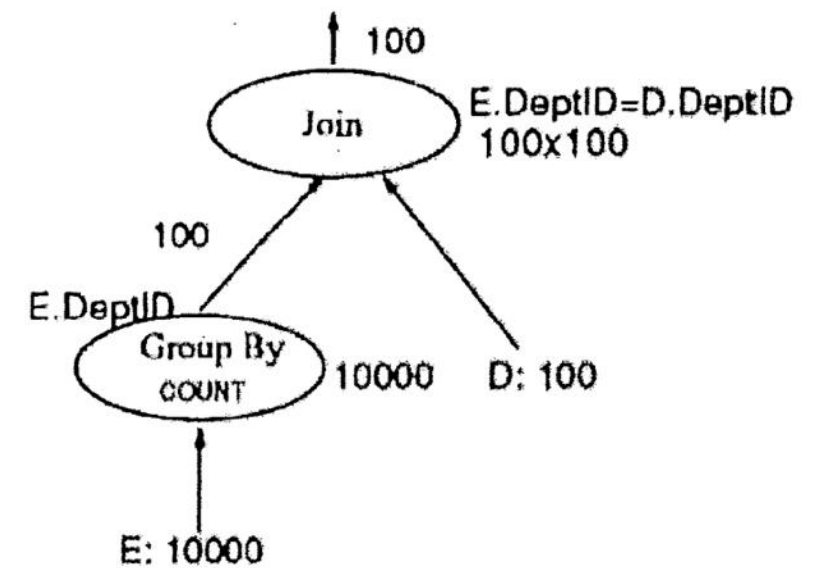
```
SELECT D.name, count(*)  
FROM EMP as E, DEPT as D  
WHERE E.DeptID = D.DeptID  
GROUP BY D.name
```

E has 10000 tuples

D has 100 tuples



Plan 1: Group by after join



Plan 2: Group by before join

Outline

Aggregation pushdown

System catalog

Cost estimation

- Cardinality estimation
- IO cost estimation

Plan enumeration

System Catalog

Catalog tables: Special tables that store the descriptive information of every table and index

Example:

attr_name	rel_name	type	position
sid	Sailors	integer	1
sname	Sailors	string	2
rating	Sailors	integer	3
age	Sailors	real	4
sid	Reserves	integer	1
bid	Reserves	integer	2
day	Reserves	dates	3
rname	Reserves	string	4

Information in the Catalog

For each table

- Table name, file name, file structure (e.g., heap file)
- Attribute name and type of each attribute
- Index name
- Integrity constraints (e.g., primary key and foreign key)

For each index

- Index name and the structure of the index
- Search key attributes

For each view

- View name and definition

Statistics About Tables and Indexes

Table cardinality: The number of tuples ***NTuples(*R*)*** for each table *R*

Statistics About Tables and Indexes

Table cardinality: The number of tuples **$NTuples(R)$** for each table R

Table size: The number of pages **$NPages(R)$** for each table R

Statistics About Tables and Indexes

Table cardinality: The number of tuples **$NTuples(R)$** for each table R

Table size: The number of pages **$NPages(R)$** for each table R

Index cardinality: The number of distinct key values **$NKeys(I)$** for each index I

Statistics About Tables and Indexes

Table cardinality: The number of tuples **$NTuples(R)$** for each table R

Table size: The number of pages **$NPages(R)$** for each table R

Index cardinality: The number of distinct key values **$NKeys(I)$** for each index I

Index size: The number of pages **$INPages(I)$** for each index I (for a B+ tree index, we take $INPages$ to be the number of leaf pages)

Statistics About Tables and Indexes

Table cardinality: The number of tuples $NTuples(R)$ for each table R

Table size: The number of pages $NPages(R)$ for each table R

Index cardinality: The number of distinct key values $NKeys(I)$ for each index I

Index size: The number of pages $INPages(I)$ for each index I (for a B+ tree index, we take $INPages$ to be the number of leaf pages)

Index height: The number of nonleaf levels $IHeight(I)$ for each index I

Statistics About Tables and Indexes

Table cardinality: The number of tuples $NTuples(R)$ for each table R

Table size: The number of pages $NPages(R)$ for each table R

Index cardinality: The number of distinct key values $NKeys(I)$ for each index I

Index size: The number of pages $INPages(I)$ for each index I (for a B+ tree index, we take $INPages$ to be the number of leaf pages)

Index height: The number of nonleaf levels $IHeight(I)$ for each index I

Index range: The minimum present key value $ILow(I)$ and the maximum present key value $IHigh(I)$ for each index I

Outline

Aggregation pushdown

System catalog

Cost estimation

- Cardinality estimation
- IO cost estimation

Plan enumeration

Cost Estimation

Estimating the cost of a query plan involves:

1. Estimating the **cost** of each operation in the plan
 - Depends on input cardinalities
 - Algorithm cost
2. Estimating the **size** of intermediate results (**cardinality estimation**)
 - We need statistics about input relations
 - For selections and joins, we typically assume independence of predicates

Outline

Aggregation pushdown

System catalog

Cost estimation

- **Cardinality estimation**

- IO cost estimation

Plan enumeration

Cardinality Estimation

Cardinality [Mathematical]: The number of values in a set

Cardinality [Database]: the number of distinct values in a table column

Cardinality estimation: Predict the number of rows that a (sub-)query returns

- Plays a crucial rule in query optimization!

Estimating Result Size

Consider the following query block

```
SELECT attribute list
FROM   relation list
WHERE  term1
      AND term2
      ...
      AND termn
```

We associate each term with a **reduction factor**: ratio of the expected result size to the input size

Estimated output size = maximum size times the product of the reduction factors for the terms in the WHERE clause

Reduction Factors

Format of terms	Reduction factor
column = value	If have index I : $\frac{1}{NKeys(I)}$ Otherwise: 1/10 (arbitrary guess)

Assume uniform distribution

If we have more statistics (i.e., number of distinct values), can do better than 1/10

Reduction Factors


Format of terms	Reduction factor
column = value	If have index I : $\frac{1}{NKeys(I)}$ Otherwise: 1/10 (arbitrary guess)
column1 = column2	If have index $I1$ and $I2$: $\frac{1}{MAX(NKeys(I1), NKeys(I2))}$

Assumes each key in the index with smaller cardinality has a matching value in the other index

Why Max(Nkeys1, Nkeys2)

Format of terms	Reduction factor
column = value	If have index I : $\frac{1}{NKeys(I)}$ Otherwise: 1/10 (arbitrary guess)
column1 = column2	If have index $I1$ and $I2$: $\frac{1}{MAX(NKeys(I1), NKeys(I2))}$

Assumes each key in the index with smaller cardinality has a matching value in the other index




ICARD1 < ICARD2

For each record in relation 1, $(NTuples2 / NKeys2)$ tuples in relation 2 will satisfy the predicate

Why Max(Nkeys1, Nkeys2)

Format of terms	Reduction factor
column = value	If have index I : $\frac{1}{NKeys(I)}$ Otherwise: 1/10 (arbitrary guess)
column1 = column2	If have index $I1$ and $I2$: $\frac{1}{MAX(NKeys(I1), NKeys(I2))}$

Assumes each key in the index with smaller cardinality has a matching value in the other index



ICARD1 < ICARD2

For each record in relation 1, $(NTuples2 / NKeys2)$ tuples in relation 2 will satisfy the predicate


Total number of selected tuples =

$$NTuples1 * NTuples2 / NKeys2$$

Why Max(Nkeys1, Nkeys2)

Format of terms	Reduction factor
column = value	If have index I : $\frac{1}{NKeys(I)}$ Otherwise: 1/10 (arbitrary guess)
column1 = column2	If have index $I1$ and $I2$: $\frac{1}{MAX(NKeys(I1), NKeys(I2))}$

Assumes each key in the index with smaller cardinality has a matching value in the other index



ICARD1 < ICARD2

For each record in relation 1, $(NTuples2 / NKeys2)$ tuples in relation 2 will satisfy the predicate

Total number of selected tuples =

$$NTuples1 * NTuples2 / NKeys2$$

Reduction factor **$F = 1 / NKeys2$**

Reduction Factors

Format of terms	Reduction factor
column = value	If have index I : $\frac{1}{NKeys(I)}$ Otherwise: 1/10 (arbitrary guess)
column1 = column2	If have index $I1$ and $I2$: $\frac{1}{MAX(NKeys(I1), NKeys(I2))}$ If have one index (e.g., $I1$): $\frac{1}{NKeys(I1)}$ Otherwise: 1/10

Reduction Factors

Format of terms	Reduction factor
column = value	If have index I : $\frac{1}{NKeys(I)}$ Otherwise: 1/10 (arbitrary guess)
column1 = column2	If have index $I1$ and $I2$: $\frac{1}{MAX(NKeys(I1), NKeys(I2))}$ If have one index (e.g., $I1$): $\frac{1}{NKeys(I1)}$ Otherwise: 1/10
column > value	If have index I : $\frac{High(I) - value}{High(I) - Low(I)}$ Otherwise: arbitrary value (e.g., 1/3)

Reduction Factors

Format of terms	Reduction factor
column = value	If have index I : $\frac{1}{NKeys(I)}$ Otherwise: 1/10 (arbitrary guess)
column1 = column2	If have index $I1$ and $I2$: $\frac{1}{MAX(NKeys(I1), NKeys(I2))}$ If have one index (e.g., $I1$): $\frac{1}{NKeys(I1)}$ Otherwise: 1/10
column > value	If have index I : $\frac{High(I) - value}{High(I) - Low(I)}$ Otherwise: arbitrary value (e.g., 1/3)
column in (list of values)	(Reduction factor of “column = value”) × (number of items in list)


Improved Statistics – Histogram

Consider **column** > **value**, We can estimate reduction factor by $\frac{High(I) - value}{High(I) - Low(I)}$, which assumes uniform distribution

Improved Statistics – Histogram

Consider **column** > **value**, We can estimate reduction factor by $\frac{High(I) - value}{High(I) - Low(I)}$, which assumes uniform distribution

We can approximate the distribution more accurately with a histogram



Outline

Aggregation pushdown

System catalog

Cost estimation

- Cardinality estimation

- **IO cost estimation**

Plan enumeration

IO cost

Calculate the number of pages access through IO

Table scan

– $IO = NPages(R)$

IO cost

Calculate the number of pages access through IO

Table scan

- $IO = NPages(R)$

Unique index matching (e.g., EMP.ID = '123')

- $IO = 1 \text{ data page} + 1-3 \text{ index page}$

IO cost

Calculate the number of pages access through IO

Table scan

- $IO = NPages(R)$

Unique index matching (e.g., EMP.ID = '123')

- $IO = 1 \text{ data page} + 1-3 \text{ index page}$

Clustered index matching

- $IO = F(preds) * (INPages(I) + NPages(R))$ # index pages & # data pages

IO cost

Calculate the number of pages access through IO

Table scan

- $IO = NPages(R)$

Unique index matching (e.g., EMP.ID = '123')

- $IO = 1 \text{ data page} + 1-3 \text{ index page}$

Clustered index matching

- $IO = F(preds) * (INPages(I) + NPages(R))$ # index pages & # data pages

Non-clustered index matching

- $IO = F(preds) * (INPages(I) + NTuples(R))$ # index pages & # data page accesses

Outline

Aggregation pushdown

System catalog

Cost estimation

- Cardinality estimation
- IO cost estimation

Query planning

Multi-Relation Query Planning

Option 1: Bottom-up optimization

- Estimate cost for individual operators (e.g., joins) and then use dynamic programming to build up the query plan

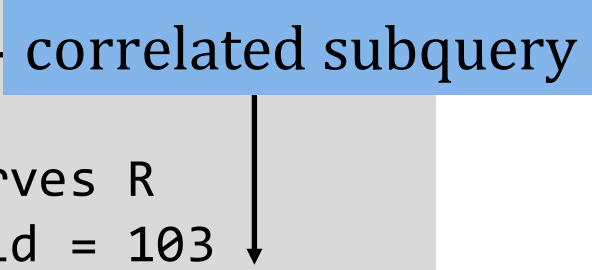
Option 2: Top-down optimization

- Start with a logical plan; search the plan tree and convert logical operators into physical operators

Nested-Query Decorrelation

Correlated
Nested query:

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT *
              FROM Reserves R
              WHERE R.bid = 103
                 AND R.sid = S.sid);
```



correlated subquery

Equivalent to:

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid = R.sid
      AND R.bid = 103;
```

Summary

Aggregation pushdown

System catalog

Cost estimation

- Cardinality estimation
- IO cost estimation

Query planning