



# CS 564: Database Management Systems

## Lecture 33: Logging

Xiangyao Yu  
4/15/2024

# Module B4 Transactions

---

Concurrency control

Optimistic concurrency control (OCC)

**Logging**

ARIES recovery

# Outline

---

## **Durability**

Types of failures

Failure examples


Write ahead logging (WAL)

Buffer management policies

Other discussion

# A Naïve Implementation of Transactions


---



**Durability:** The database must recover to a valid state no matter when a crash occurs


- Committed transactions should persist
- Uncommitted transactions should roll back

# A Naïve Implementation of Transactions




**Durability:** The database must recover to a valid state no matter when a crash occurs

- Committed transactions should persist
- Uncommitted transactions should roll back




# A Naïve Implementation of Transactions



**Durability:** The database must recover to a valid state no matter when a crash occurs

- Committed transactions should persist
- Uncommitted transactions should roll back



T1 and T2 commit;  
T3, T4 should be aborted after restart

# Outline

---

Durability

**Types of failures**

Failure examples

Write ahead logging (WAL)

Buffer management policies

Other discussion

# Types of Failures

---

## Software failure

- Crash of the operating system, the DBMS, or the application software

## Hardware failure

- Crashes due to memory or disk errors
- Power failures
- Storage media failure



# Types of Failures

---

## Software failure

- Crash of the operating system, the DBMS, or the application software

## Hardware failure

- Crashes due to memory or disk errors
- Power failures
- Storage media failure

**Logging can handle these failures**

# Types of Failures

---

## Software failure

- Crash of the operating system, the DBMS, or the application software

## Hardware failure

- Crashes due to memory or disk errors
- Power failures
- Storage media failure

**Requires data replication, beyond the scope of this class**

# Outline

---

Durability

Types of failures

**Failure examples**

Write ahead logging (WAL)

Buffer management policies


Other discussion

# Failure Example 1

---

On a crash, data in disk persists, data in memory disappears

- Failure recovery depends on buffer management policies




# Failure Example 1

---


On a crash, data in disk persists, data in memory disappears

- Failure recovery depends on buffer management policies




# Failure Example 1

- On a crash, data in disk persists, data in memory disappears
- Failure recovery depends on buffer management policies



# Failure Example 1

- On a crash, data in disk persists, data in memory disappears
- Failure recovery depends on buffer management policies

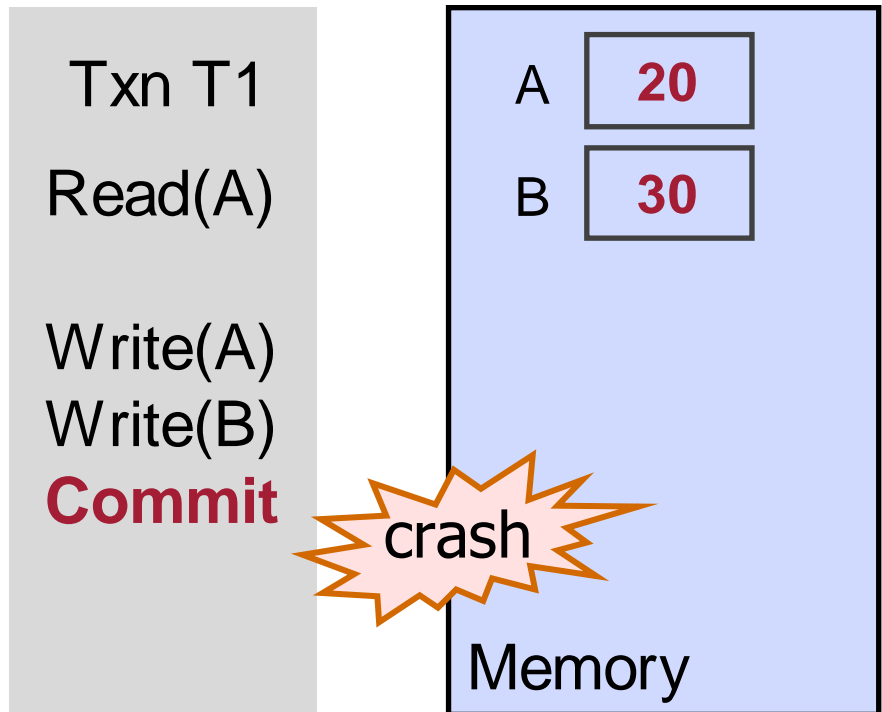


Update not reflected on disk, no need to recovery this transaction


# Failure Example 2

On a crash, data in disk persists, data in memory disappears

- Failure recovery depends on buffer management policies



Disk not updated before transaction commits






# Failure Example 2

On a crash, data in disk persists, data in memory disappears

- Failure recovery depends on buffer management policies




For recoverability, updates must be stored on disk before transaction commits

# Failure Example 3

On a crash, data in disk persists, data in memory disappears


- Failure recovery depends on buffer management policies



# Failure Example 3

On a crash, data in disk persists, data in memory disappears

- Failure recovery depends on buffer management policies



The transaction does not commit before failure, must rollback updates on disk

# Outline

---

Durability

Types of failures

Failure examples

**Write ahead logging (WAL)**

Buffer management policies

Other discussion

# Write-Ahead Logging (WAL)

---

## Write-ahead logging

- Flush a **log record** before updating the data page on disk
- Log incurs only sequential IO

# Write-Ahead Logging (WAL)

---

## Write-ahead logging

- Flush a **log record** before updating the data page on disk
- Log incurs only sequential IO


**REDO information:** information about the change made by the transaction (e.g., new value in the page)

**UNDO information:** information to reverse a change by the transaction (e.g., old value in the page)

# Write-Ahead Logging (WAL)

## Write-ahead logging

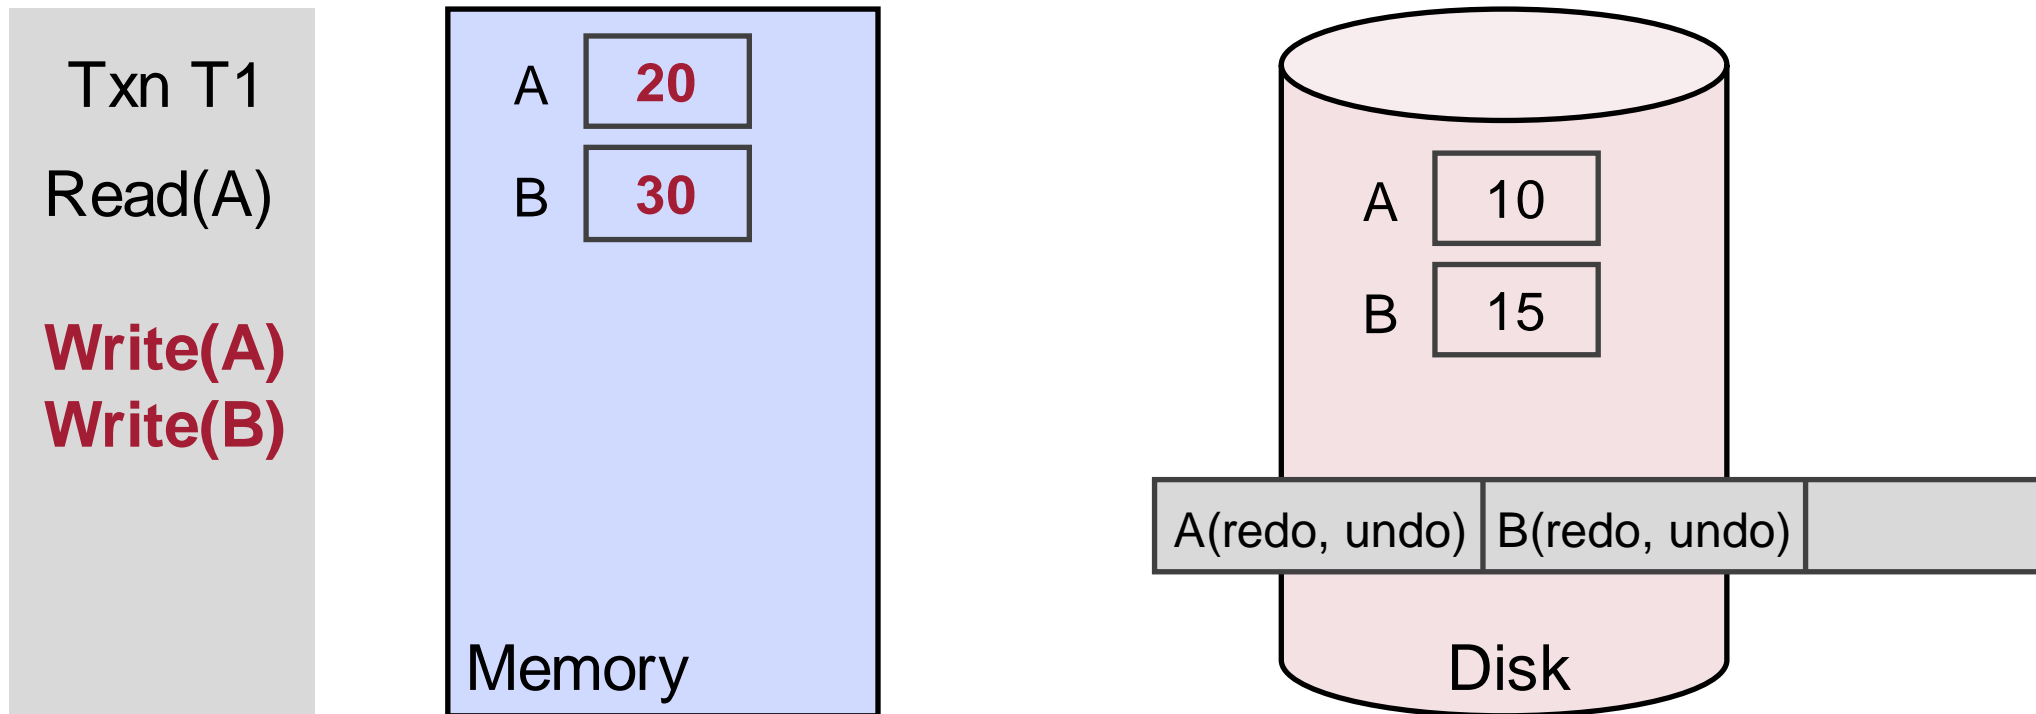
- Flush a **log record** before updating the data page on disk



# Write-Ahead Logging (WAL)

## Write-ahead logging

- Flush a **log record** before updating the data page on disk



Use REDO info to reapply changes; use UNDO info to reverse changes



# Outline

---

Durability

Types of failures

Failure examples

Write ahead logging (WAL)

**Buffer management policies**

- Steal vs. No Steal
- Force vs. No Force

Other discussion

# Buffer Management Policy

---

**No Steal:** Dirty pages stay in DRAM until the transaction commits

**Steal:** Dirty pages can be flushed to disk before the transaction commits

Why called steal? Another transaction can “steal” the buffer slot that the current transaction is using and thus the current transaction will flush the dirty page to disk

# Buffer Management Policy

---

**No Steal:** Dirty pages stay in DRAM until the transaction commits

- No need to reverse changes during recovery
- No need to log UNDO information

**Steal:** Dirty pages can be flushed to disk before the transaction commits

# Buffer Management Policy

---

**No Steal:** Dirty pages stay in DRAM until the transaction commits

- No need to reverse changes during recovery
- No need to log UNDO information

**Steal:** Dirty pages can be flushed to disk before the transaction commits

- Log UNDO information so that dirty changes can be reversed during recovery
- **Advantage:** other transactions can use the buffer slot in DRAM
- **Disadvantage:** need to log UNDO information

# Buffer Management Policy

---

**Force:** All dirty pages must be flushed before the transaction commits

**No Force:** Dirty pages may stay in memory after the transaction commits

# Buffer Management Policy

---

**Force:** All dirty pages must be flushed before the transaction commits

- No need to replay changes during recovery
- No need to log REDO information

**No Force:** Dirty pages may stay in memory after the transaction commits

# Buffer Management Policy

---

**Force:** All dirty pages must be flushed before the transaction commits

- No need to replay changes during recovery
- No need to log REDO information

**No Force:** Dirty pages may stay in memory after the transaction commits

- Log REDO information so that committed changes can be replayed during recovery
- **Advantage:** can reduce random disk IO
- **Disadvantage:** need to log REDO information

# Buffer Management Policy — Summary

---

	Steal	No Steal
Force		
No Force		<b>REDO</b> only



# Buffer Management Policy

---

	Steal	No Steal
Force		
No Force		<b>REDO</b> only

REDO only (e.g., main memory database)

- Log **REDO** record before each update to database
- Log **COMMIT** record when a transaction finishes execution
- Modified data pages can be written to disk **after COMMIT record is logged**
- **Recovery**: **replay REDO records** for committed transactions

# Buffer Management Policy

---

	Steal	No Steal
Force	<b>UNDO</b> only	
No Force		<b>REDO</b> only

## UNDO only (e.g., NVM-based database)

- Log **UNDO** record before each update to database
- Log **COMMIT** record when a transaction finishes execution
- Modified data pages can be written to disk **after the corresponding UNDO record but before the COMMIT record**
- **Recovery:** **UNDO** uncommitted transactions

# Buffer Management Policy

---

	Steal	No Steal
Force	<b>UNDO</b> only	No REDO nor UNDO
No Force		<b>REDO</b> only

## No REDO and no UNDO

- **Atomically** write all updates to disk
- Check out [1] if you are interested in more details

# Buffer Management Policy

	Steal	No Steal
Force	<b>UNDO</b> only	No REDO nor UNDO
No Force	<b>REDO and UNDO</b> logging (ARIES [2])	<b>REDO</b> only

## REDO and UNDO (the most flexible!)

- Modified data pages can be written to disk **any time after the corresponding log record**
- Log both **REDO and UNDO** before updating data pages
- Log **COMMIT** record when a transaction finishes execution
- **Recovery**: **UNDO uncommitted** transactions and **REDO committed** transactions

# Outline

---

Durability

Types of failures

Failure examples

Write ahead logging (WAL)

Buffer management policies

## **Other discussion**

- Checkpoints
- Group commit
- Physical vs. logical logging

# Checkpoints

---

The log file can grow unbounded over time

Solution: **checkpoint**

Naïve implementation of a checkpoint

- Stop admitting new transactions and commit/abort all ongoing transactions
- Flush all dirty pages to disk
- Log CHECKPOINT (Can truncate the log at this point)
- Resume execution

# Group Commit

---

Flushing a log record to disk for each write incurs significant performance overhead

Observation: Database log is a sequential structure that captures log records from all transactions

**Group commit** flushes multiple log records with a single disk IO

- **Tradeoff**: each write has a slightly higher latency, but the logging throughput is significantly improved

# Physical vs. Logical Logging

---

## Physical logging

- Record the byte changes
- E.g., REDO contains the value in the new page; UNDO contains the value in the old page

## Logical logging

- Record the logical operations (UPDATE, INSERT, DELETE)
- **Advantages:** consumes less disk space than physical logging

## Physiological logging

- Physical to a page, and logical within a page



# Summary

---

Durability

Types of failures

Failure examples

Write ahead logging (WAL)

Buffer management policies

- Steal vs. No Steal
- Force vs. No Force

Other discussion

- Checkpoints
- Group commit
- Physical vs. logical logging